

<https://tinyurl.com/slack-nuces>

# Operating Systems Design

## 3. Definitions, Concepts, and Architecture

Paul Krzyzanowski  
pxk@cs.rutgers.edu

How      What

Mechanisms & Policies

# OS Mechanisms & Policies

---

- Mechanisms:
  - Presentation of a software abstraction:
    - Memory, data blocks, network access, processes
- Policies:
  - Procedures that define the behavior of the mechanism
    - Allocation of memory regions, replacement policy of data blocks
  - Permissions
- Keep mechanisms, policies, and permissions separate

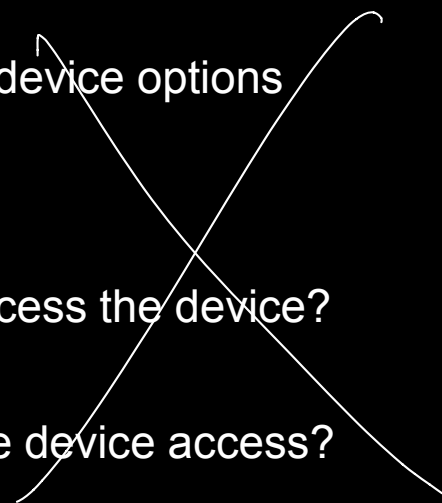
# Processes

---

- Mechanism:
  - Create, terminate, suspend, switch, communicate
- Policy
  - Who is allowed to create and destroy processes?
  - What is the limit?
  - What processes can communicate?
  - Who gets priority?

# Character Devices

---

- Mechanism:
    - Read, write, change device options
  - Policy
    - Who is allowed to access the device?
    - Is sharing permitted?
    - How do you schedule device access?
- 

# Definitions, Concepts, and Architecture

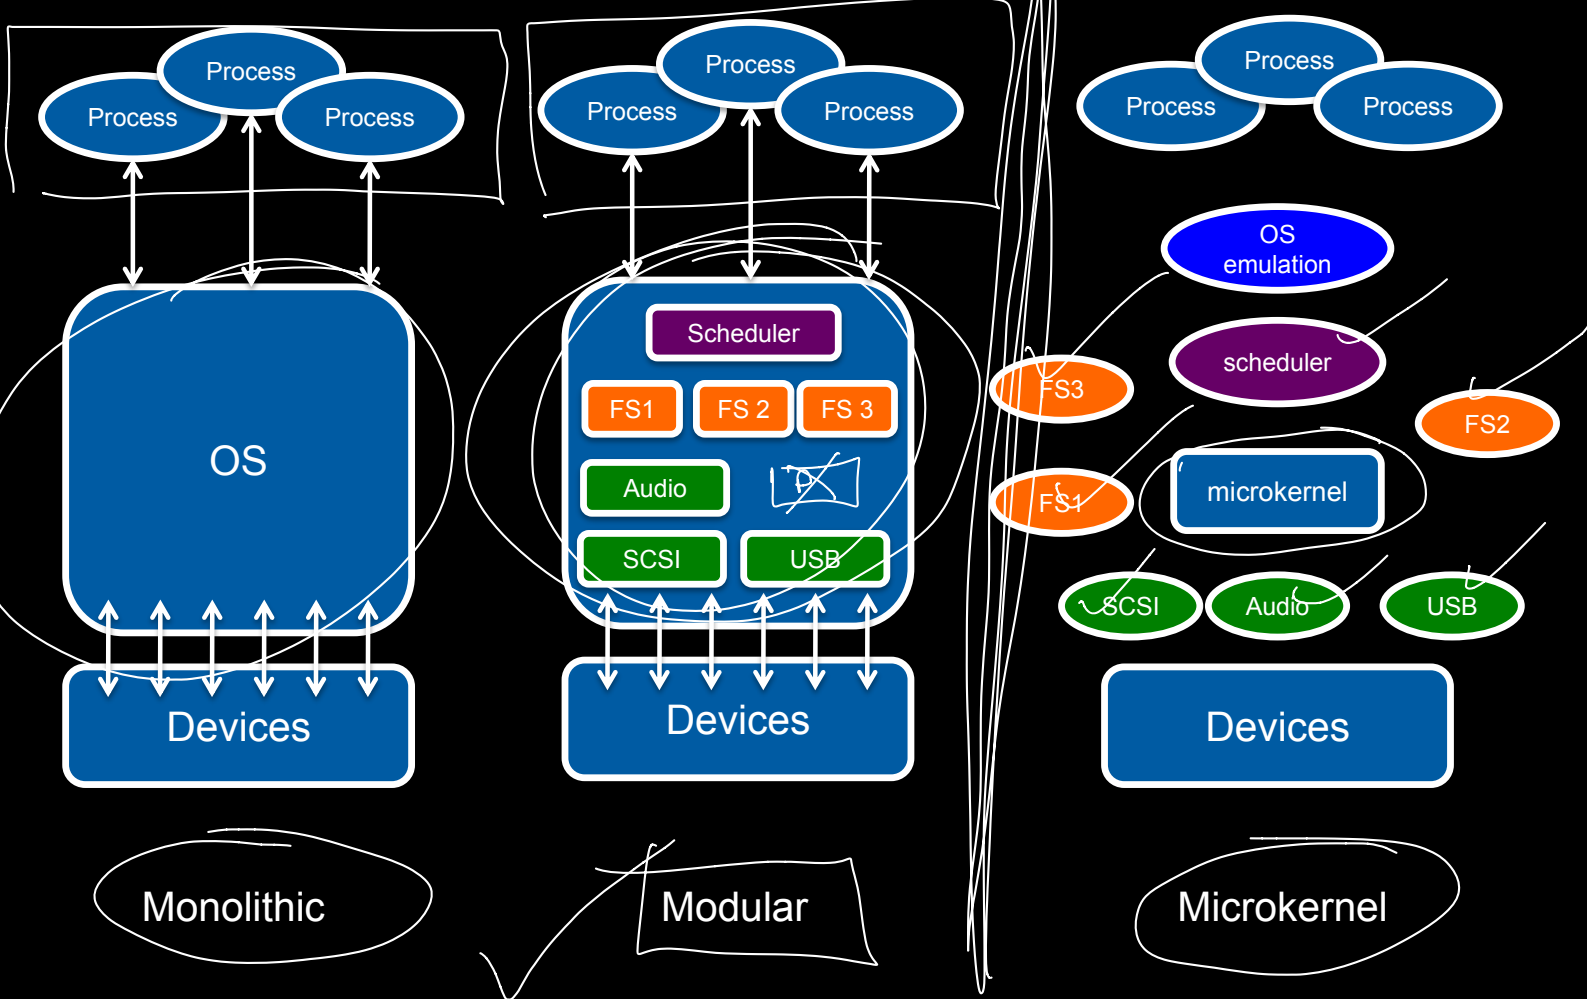
# What is an operating system?

---

- The first program
- A program that lets you run other programs
- A program that provides controlled access to resources:
  - CPU
  - Memory
  - Display, keyboard, mouse
  - Persistent storage
  - Network

This includes: naming, sharing, protection, communication

# OS Structure





# What's a kernel?

---

- **Operating System**

- Often refers to the complete system, including command interpreters, utility programs, window managers, ...

- **Kernel**

- Core component of the system that manages resource access, memory, and process scheduling

# UNIX Kernel (example)

---

## Some of the things it does:

- Controls execution of processes
  - Creation, termination, communication
- Schedules processes for execution on the CPU(s)
- Manages memory
  - Allocates memory for an executing process
  - Sets memory protection
  - Coordinates swapping pages of memory to a disk if low on memory
- Manages a file system
  - Allocation and retrieval of disk data
  - Enforcing access permissions & mutual exclusion
- Provides access to devices
  - Disk drives, networks, keyboards, displays, printers, ...
  - Enforces access permissions & mutual exclusion

# User Mode vs. Kernel Mode

- **Kernel mode** = privileged, system, supervisor mode

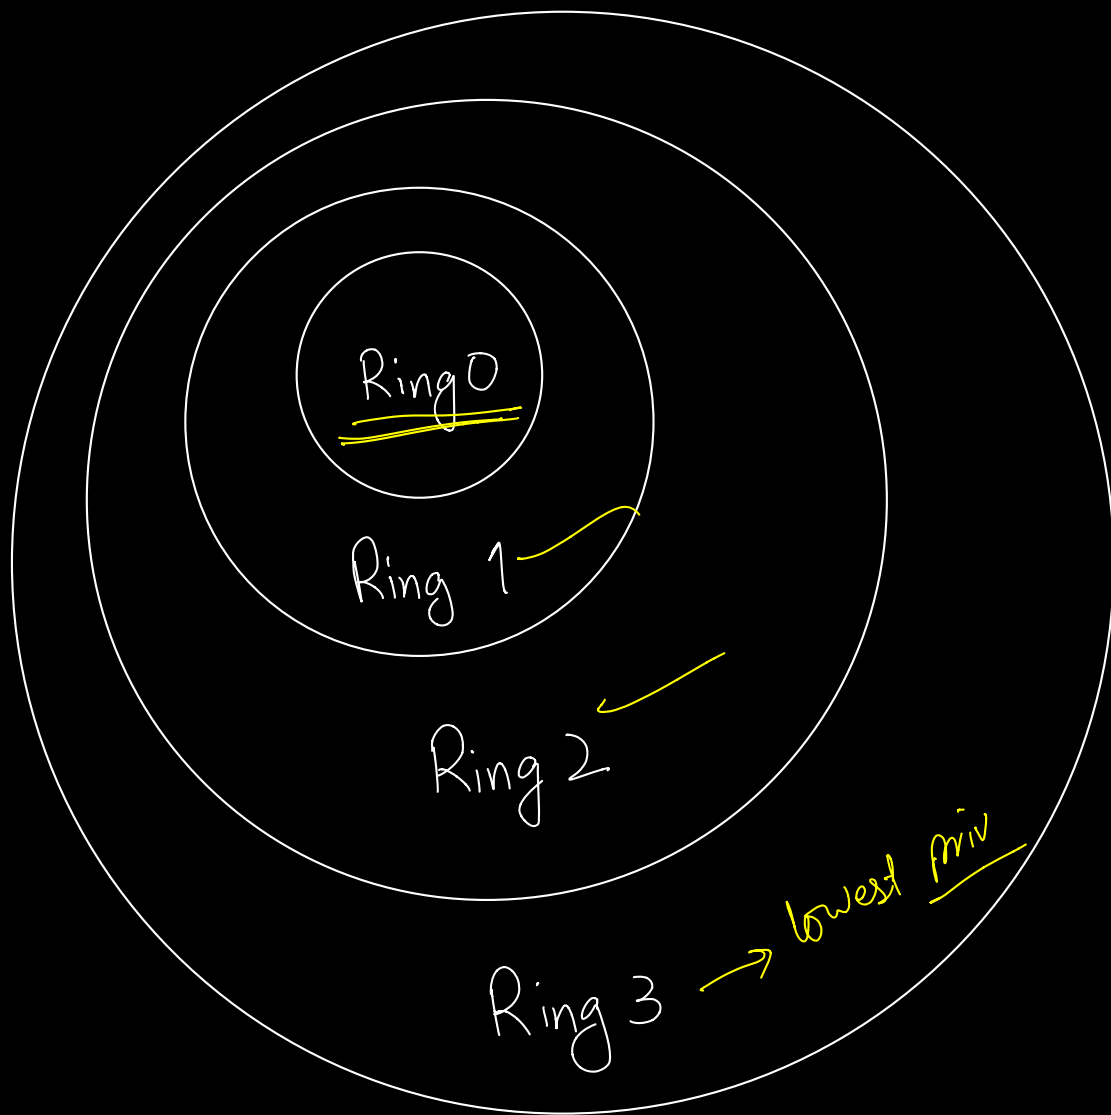
- Access restricted regions of memory
- Modify the memory management unit
- Set timers
- Define interrupt vectors
- Halt the processor
- Etc.

*can do anything*

- CPU knows what mode it's in via a status register

- You can set the register in kernel mode
- OS & boot loaders run in kernel mode
- User programs run in user mode





User mode  
(userspace)

---

Kernel mode  
(kernel space)

# Violations

- What if a CPU tries to execute something that is available only in kernel mode?
  - (a) nothing or (more likely)
  - (b) trap (exception)
    - Memory access violation
    - Illegal instruction violation
    - Register access violation
- The OS processes the trap
  - Original program counter is saved
  - OS decides on course of action
  - If needed, restart the offending instruction
- Traps occur:
  - Via software (e.g., INT instruction)
  - Because of an access violation
  - Via a hardware interrupt (e.g., timer)

# How do you switch to kernel mode?


## Software interrupts (traps)

- Trap vectors are set up in kernel mode (at boot time)
  - Trap pushes the return address on the stack and jumps to a well-known address
  - That address usually contains a *jump* instruction (vector) to the code that will handle that trap
- Returning back to user mode: *return from exception*

**Mode Switch:** switching between user & kernel mode

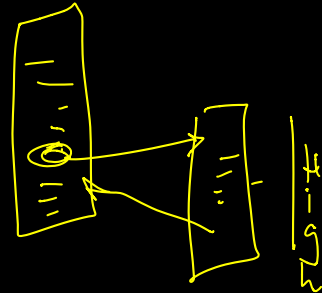
EAX → (10)  
INT → 80h

~~INT~~



# System Calls: Interacting with the OS

- Use *trap* mechanism to switch to the kernel
  - Mode switch
- Pass a number that represents the OS service
  - System call number; usually set in a register
- A system call involves:
  - Set system call number
  - Save parameters
  - Issue the trap (jump to kernel mode)
    - OS gets control
    - Return from exception (back to user mode)
  - Retrieve results and return them to the calling function
- System call interfaces are encapsulated as library functions



# Interrupts & Preemption

INT / HW

- How do we ensure that the OS gets control?
- Program a timer interrupt
  - On Linux/Intel systems,  
Set the 8254 Programmable Interval Timer to generate an interrupt (IRQ 0) approximately every 10 ms.
  - Since 2005: High Precision Event Timer (HPET) replaces 8254

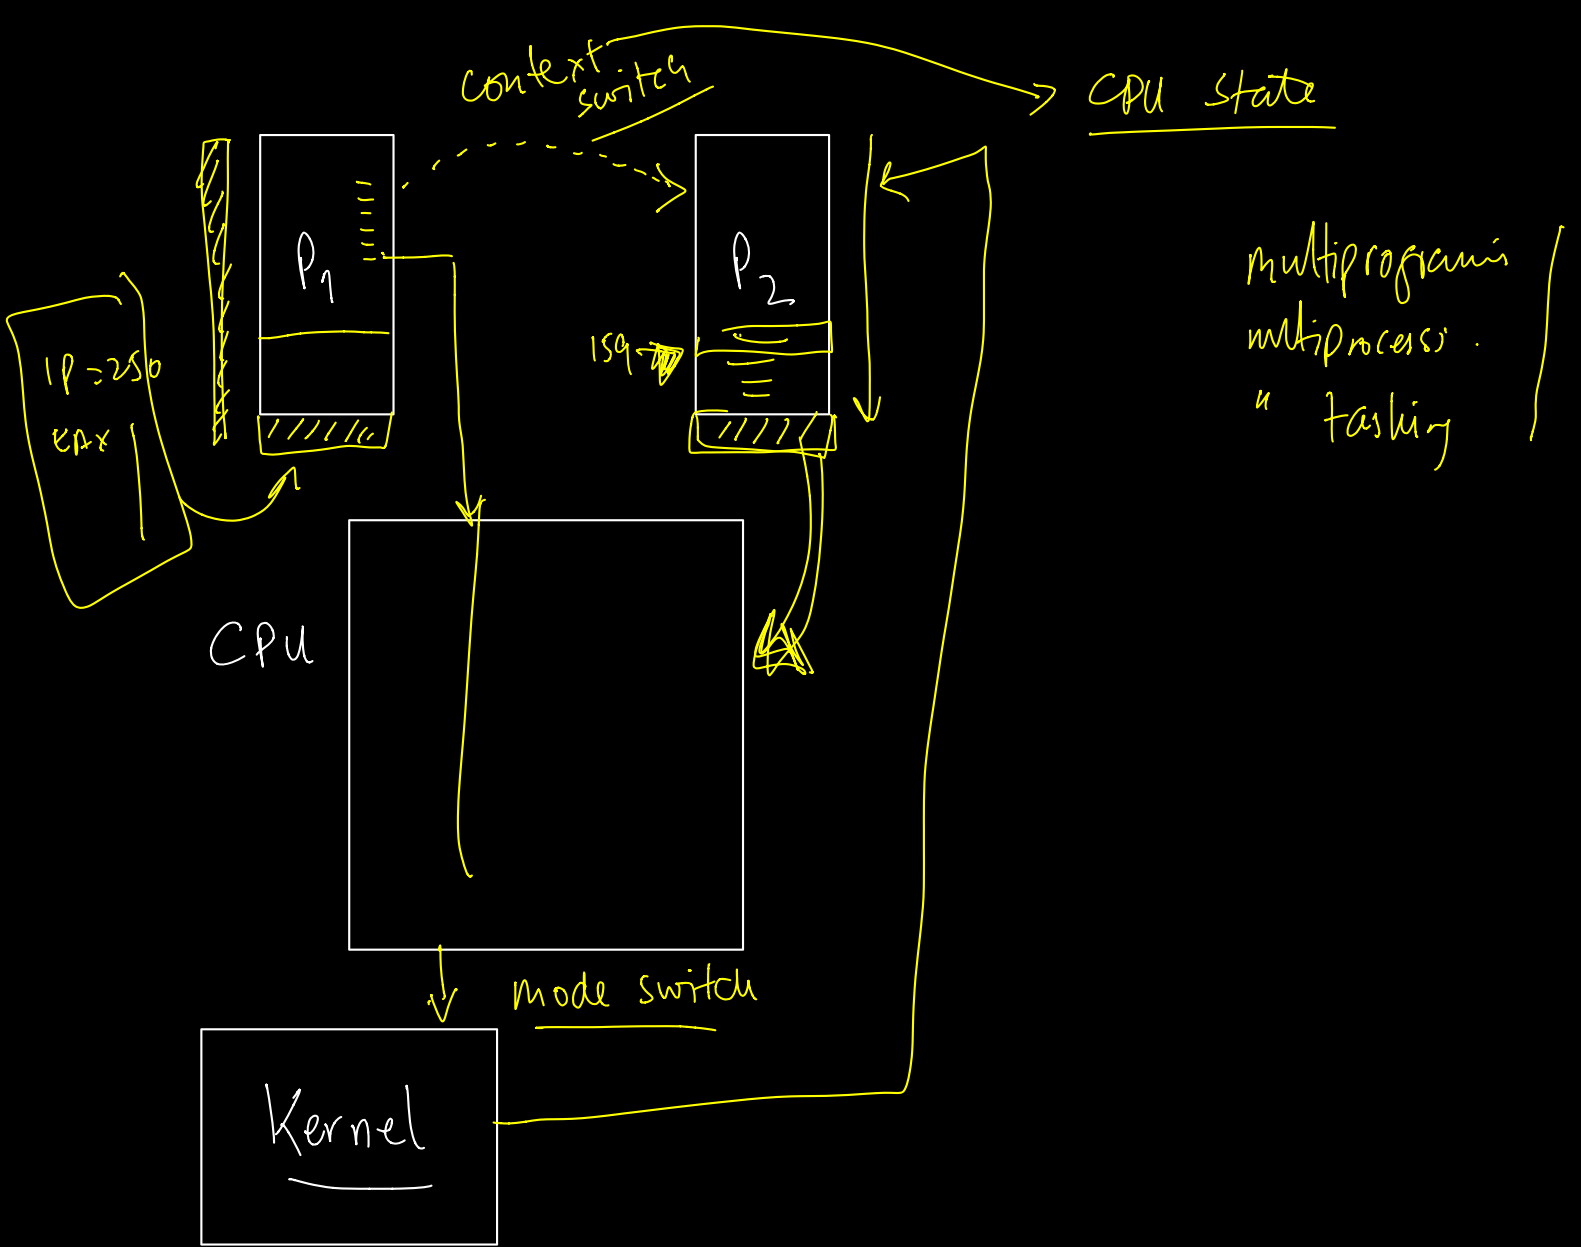
(yield X)



# Context switch & Mode switch

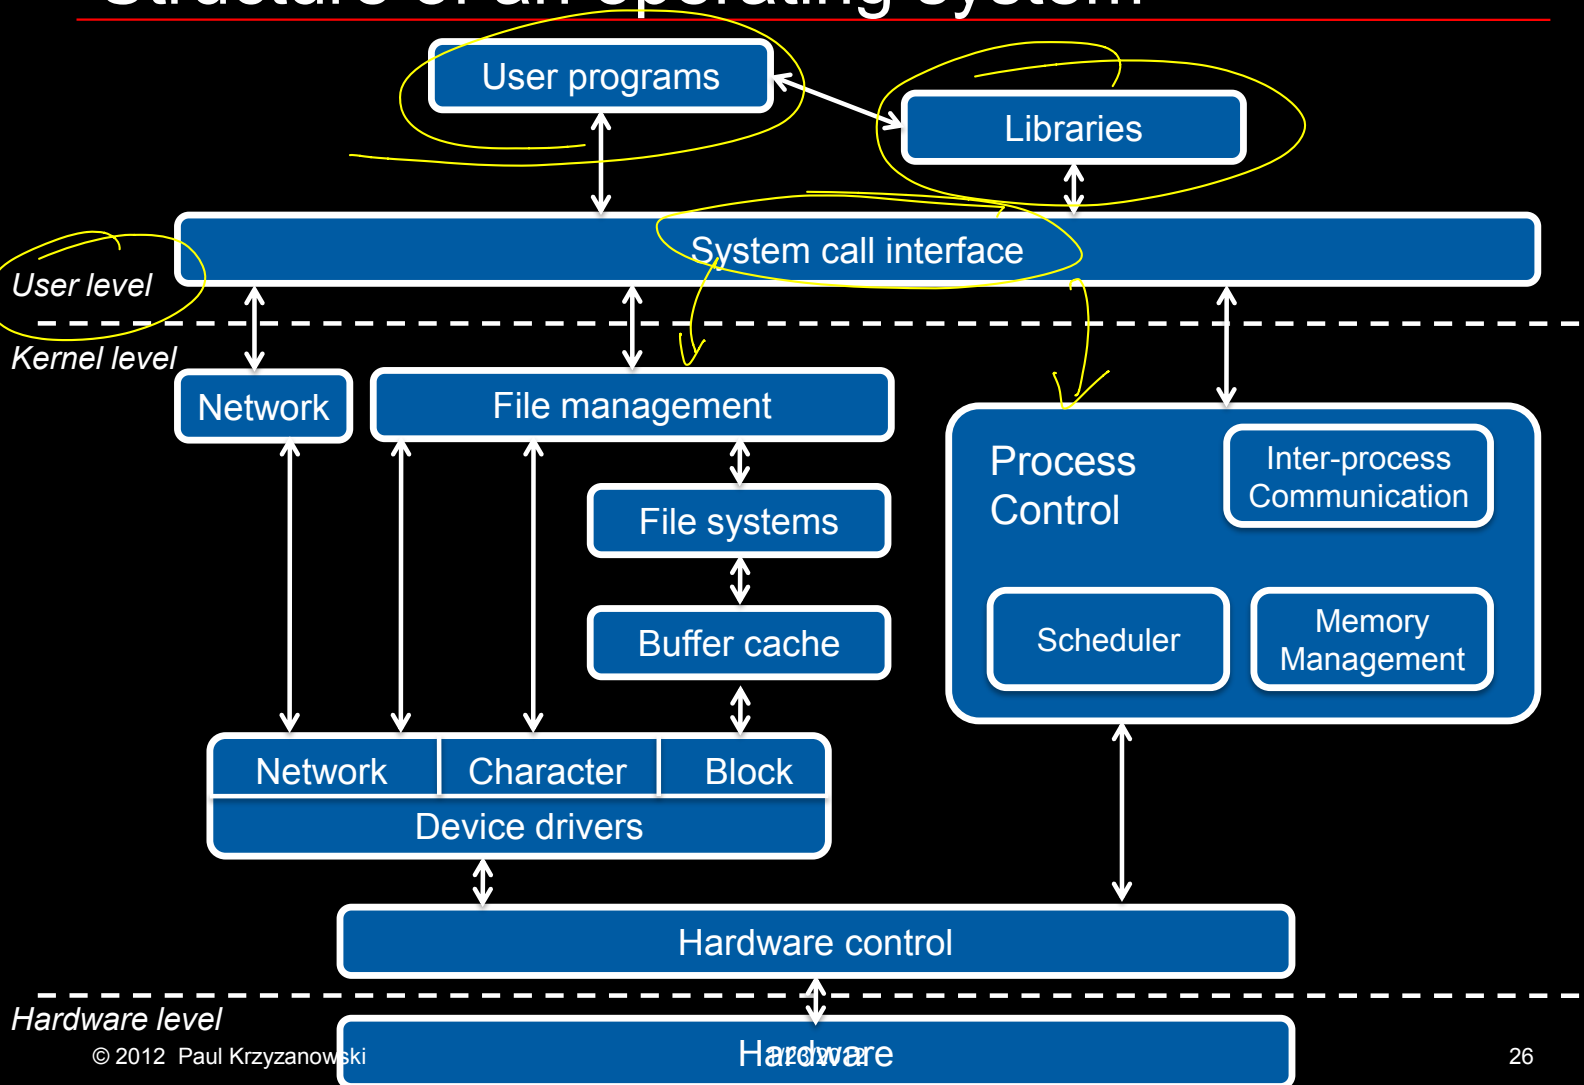
---

- An interrupt or trap results in a *mode switch*
  - CPU switches execution from user mode to kernel mode
- An operating system may save a process' state and restore another process' state.
  - *Context switch*
  - Save all registers  
(including stack pointers, PC, and flags)
  - Load saved registers (including SP, PC, flags)
  - To return to original context: restore registers and return from exception
- Context switch: switch to kernel mode, save state so that it can be restored later and reload another process' saved state



— system calls.

# Structure of an operating system



# POSIX

*Syscalls*

*Win32 API*

- UNIX → POSIX
- IEEE (ISO/IEC 9945): defines POSIX environment
  - System interfaces
  - Shell & scripting interface
  - Common utilities
  - Networking interfaces
  - Security interfaces
- POSIX (or close to) systems include
  - Solaris, BSD, Mac OS X, VxWorks, Microsoft Windows Services for UNIX
  - Linux, FreeBSD, NetBSD, OpenBSD, BeOS

OSDT

INT

SYSENTER

|||

The End.