

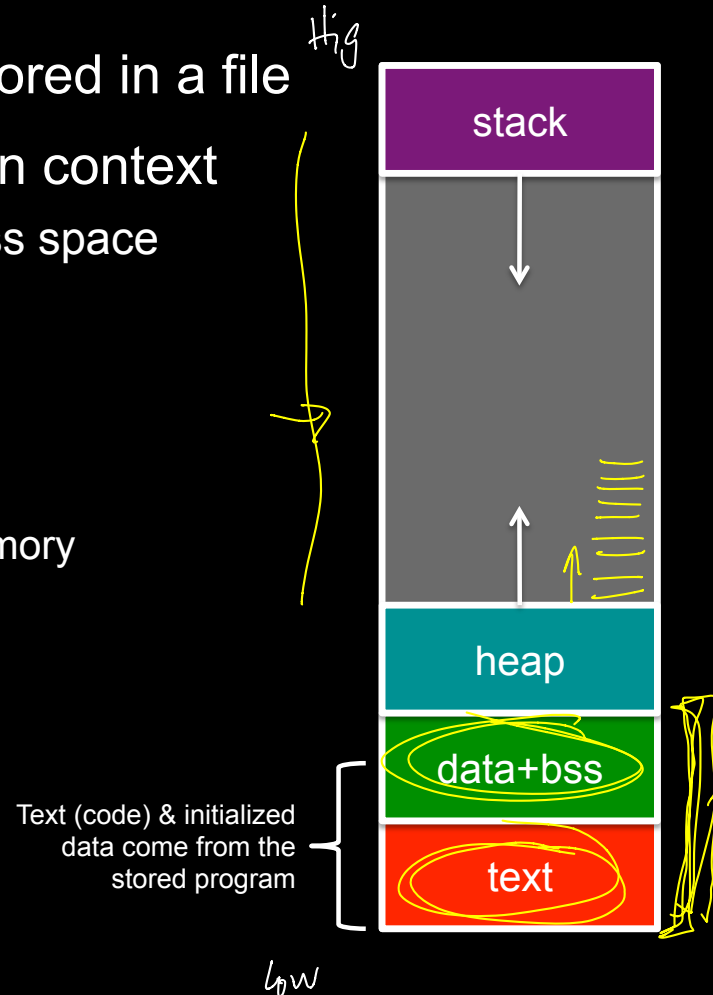
Operating Systems Design

4. Processes

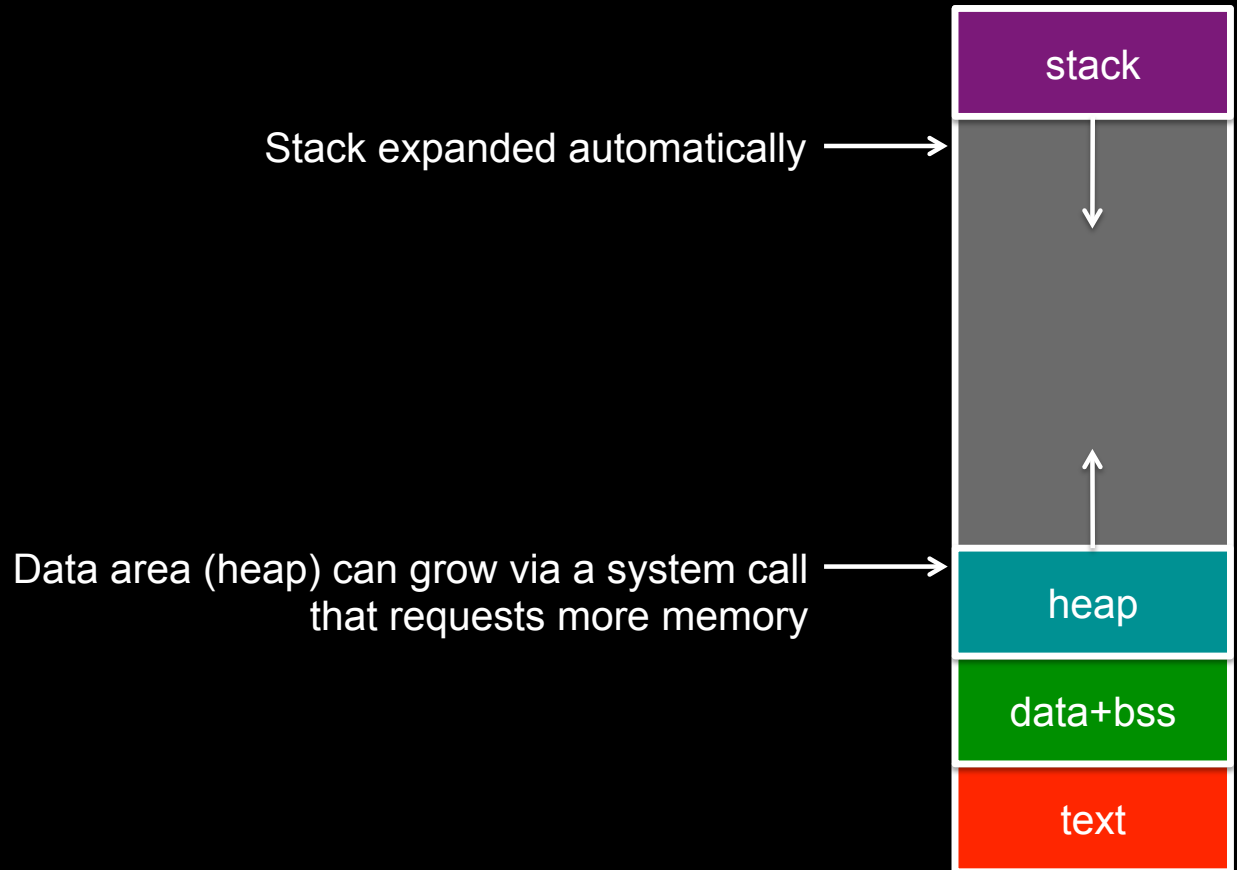
Paul Krzyzanowski
pxk@cs.rutgers.edu

Process

- **Program**: code & static data stored in a file
- **Process**: a program's execution context
 - Each process has its own address space
 - **Memory map** *mapping*
 - **Text**: compiled program
 - **Data**: initialized static data
 - **BSS**: uninitialized static data
 - **Heap**: dynamically allocated memory
 - **Stack**: call stack
 - **Process context**:
 - Program counter
 - CPU registers



Growing memory



Contexts

- Entering the kernel
 - Hardware interrupts
 - Asynchronous events (I/O, clock, etc.)
 - Do not relate to the context of the current process
 - Because they are asynchronous, *any* process might be running when they occur
 - Software traps
 - Are related to the context of the current process [process context]
 - Examples: illegal memory access, divide by zero, illegal instruction
 - Software initiated traps
 - System call from the current process [process context]
 - *The current executing process' address space is active on a trap*
- Saving state
 - Kernel stack switched in upon entering kernel mode
 - Kernel must save machine state before servicing event
 - Registers, flags (program status word), program counter, ...

read

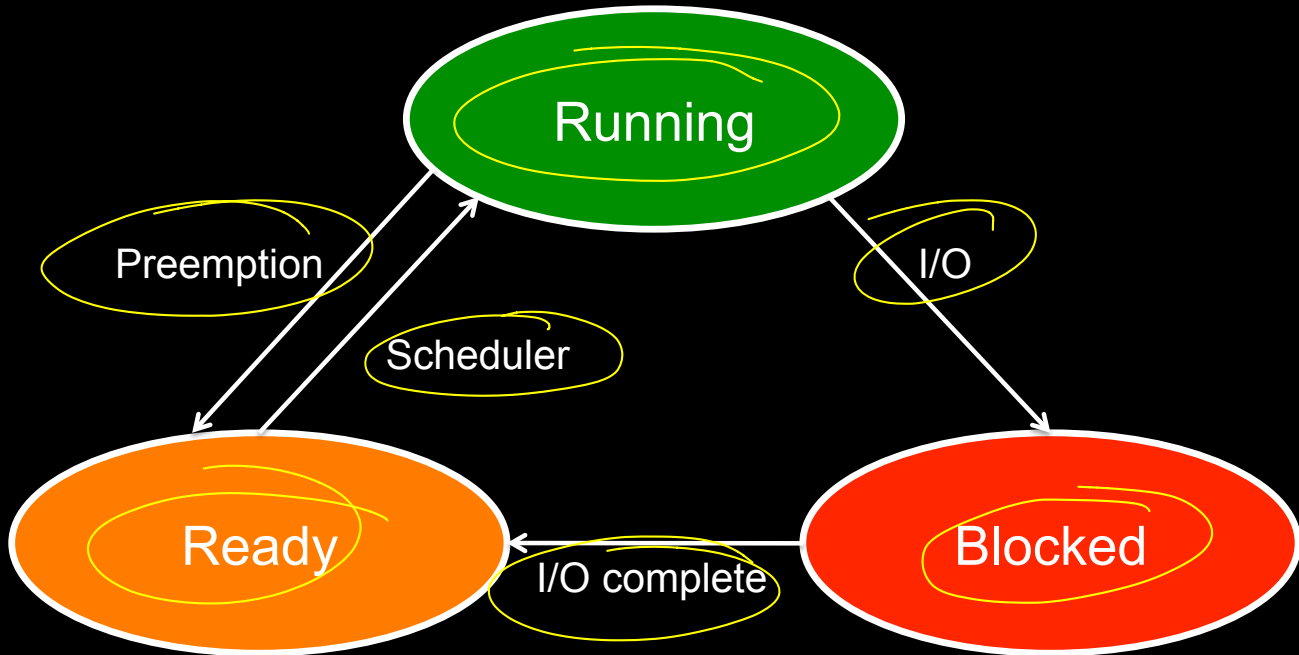
System calls

- **Entry: Trap to system call handler**
 - Save state
 - Verify parameters are in a valid address
 - Copy them to kernel address space
 - Call the function that implements the system call
 - If the function has to (cannot be satisfied immediately) then
 - Context switch to let another ready process run
 - Put our process on a blocked list
- **Return from system call or interrupt**
 - Check for signals to the process
 - Call the appropriate handler if signal is not ignored
 - Check if another process should run
 - Context switch to let the other process run
 - Put our process on a *ready* list
 - Calculate time spent in the call for profiling/accounting
 - Restore user process state
 - Return from interrupt

Processes in a Multitasking Environment

- Multiple concurrent processes
 - Each has a unique identifier: **Process ID (PID)**
- Asynchronous events (interrupts) may occur
- Processes may request operations that take a long time
- Goal: have some process running at all times
- Context saving/switching
 - Processes may be suspended and resumed
 - Need to save all state about a process so we can restore it

Process states



CPU-bound process \longrightarrow (Almost) no I/O


Keeping track of processes

- Process list stores a **Process Control Block (PCB)** per process
- A Process Control Block contains: *Structure*
 - Process ID
 - Machine state (registers, program counter, stack pointer)
 - Parent & list of children
 - Process state (ready, running, blocked)
 - Memory map
 - Open file descriptors
 - Owner (user ID) – determine access & signaling privileges
 - Event descriptor if the process is blocked
 - Signals that have not yet been handled
 - Policy items: Scheduling parameters, memory limits
 - Timers for accounting (time & resource utilization)
 - (Process group)

Processes in Linux

- The OS creates one task on startup:
 - init*: the parent of all tasks
 - launchd*: replacement for *init* on Mac OS X and FreeBSD
- Process state stored in `struct task_struct`
 - Defined in `linux/sched.h`
- Stored as a circular, doubly linked list
 - `struct list_head` in `linux/list.h`

```
struct task_struct init_task; /* static definition */
```



Aside: The Linux kernel

① - Download from: `kernel.org` (latest stable)

② - Unzip file: `unxz filename.tar.xz`

③ - Extract files: `tar xf filename.tar`

④ - Open ^{any} file:
e.g `vi include/linux/sched.h`

(We will build the kernel soon inshaallah)

`grep -r " " .`
grep -r " " .
whole tree any string dot (starting location)


Processes in Linux

- Iterating through processes

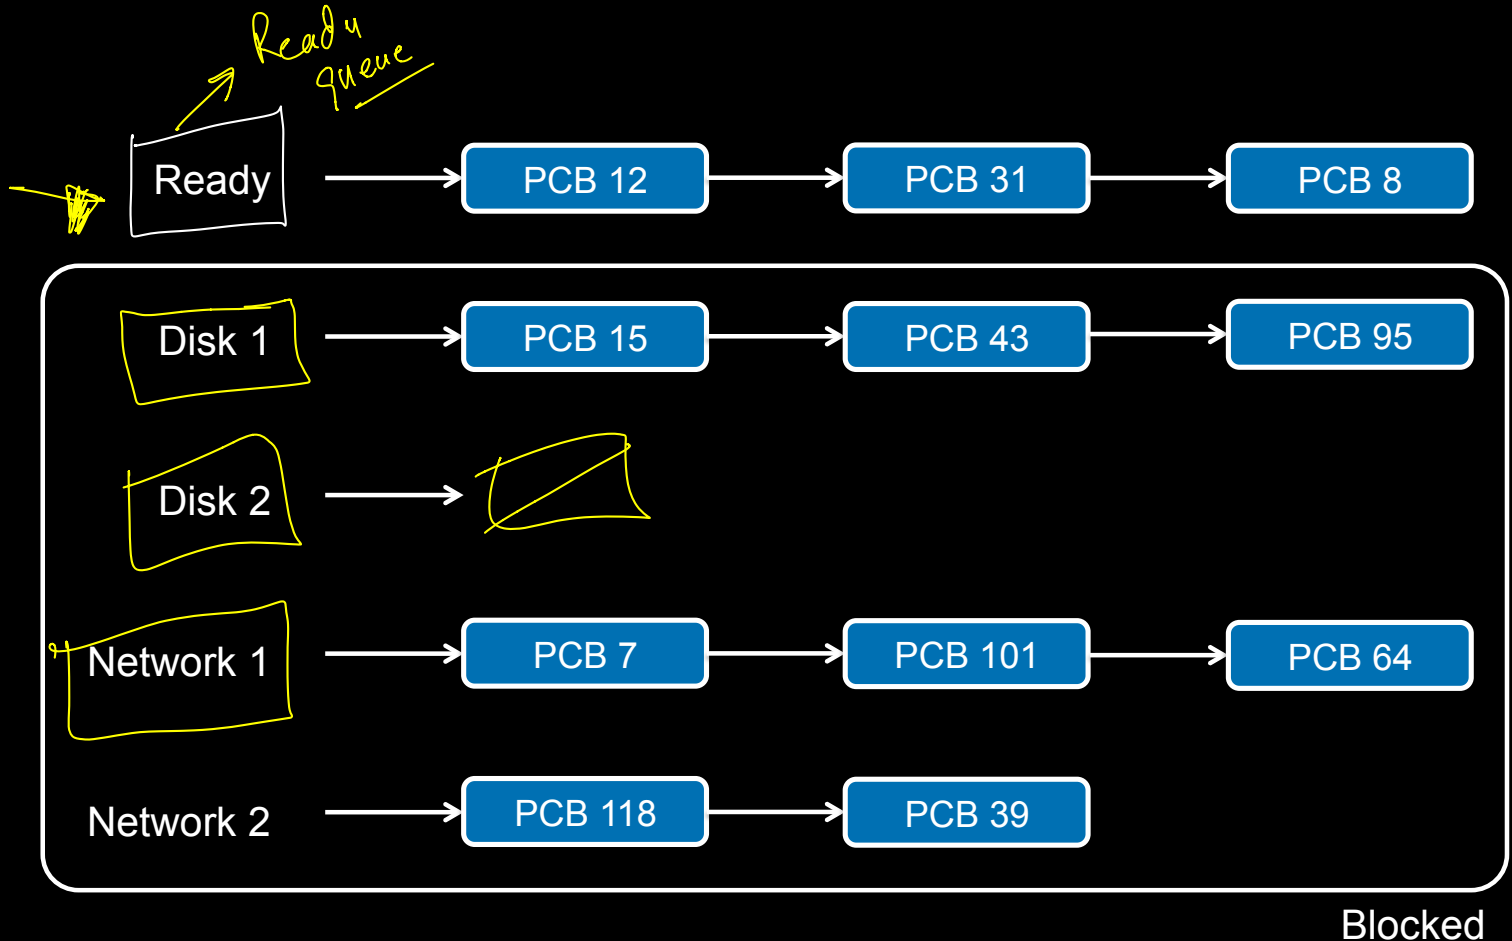
```
for (p = &init_task; ((p = next_task(p)) != &init_task; ) {  
    /* whatever */  
}
```

- The current process on the current CPU is obtained from the macro `current`

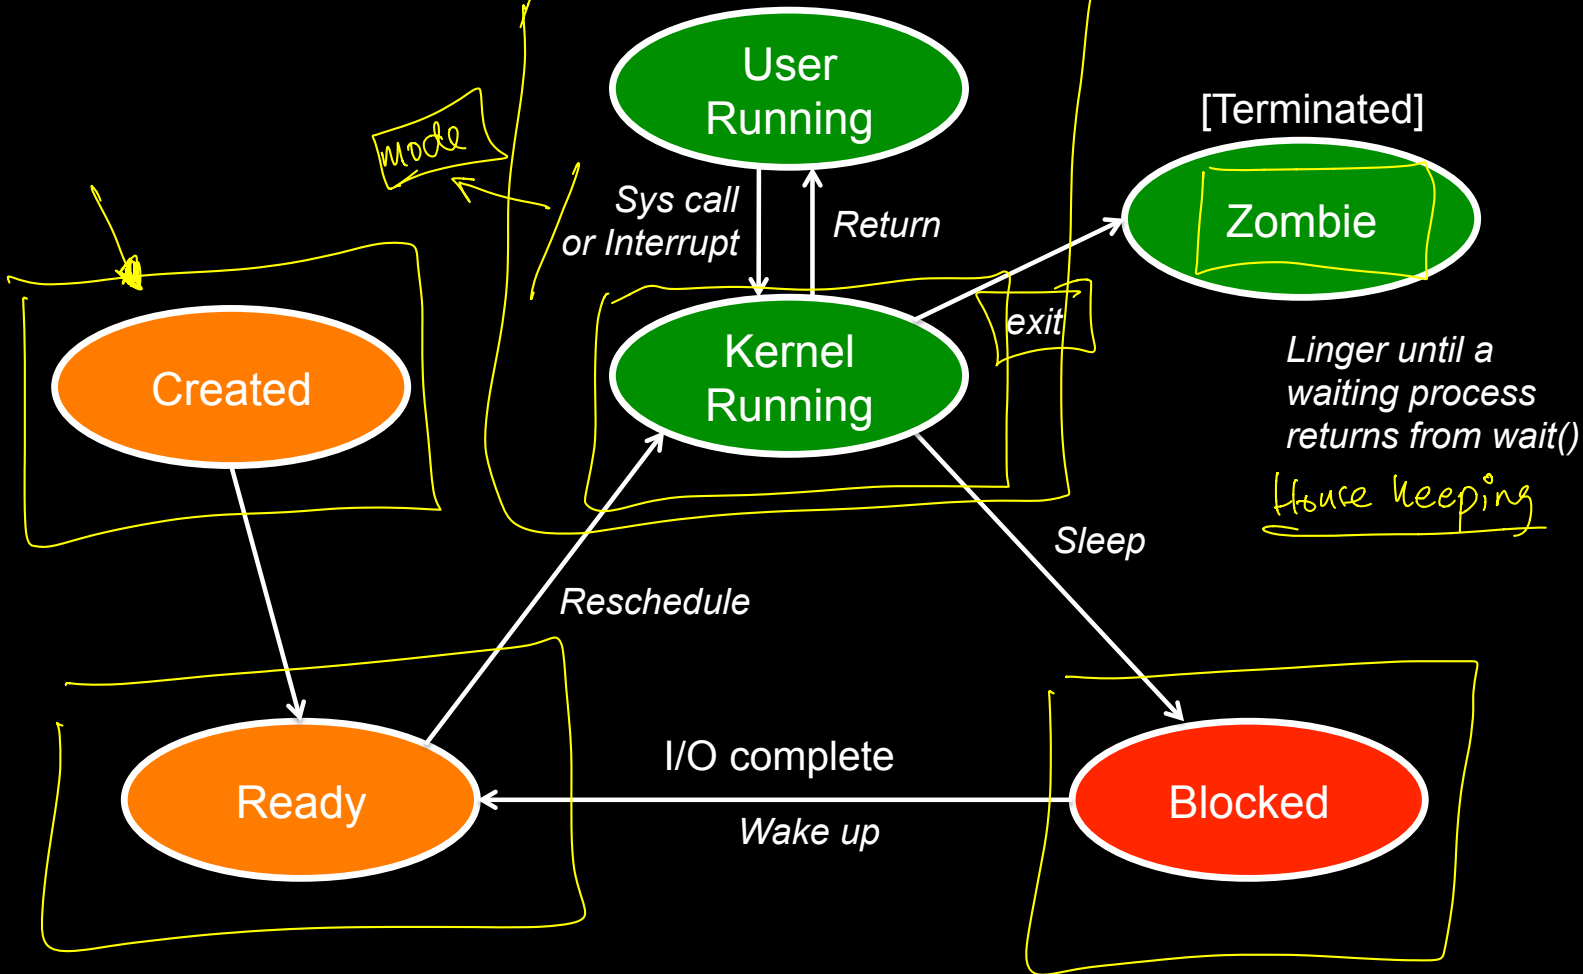
```
current->state = TASK_STOPPED;
```



Processes on Ready & Blocked Queues



Process States: a bit more detail



Creating a process under POSIX

fork system call


- Clones a process into two processes
 - New context is created: duplicate of parent process
- *fork* returns 0 to the child and the process ID to the parent

"fork in the road"

What happens?

- Check for available resources
- Allocate a new PCB
- Assign a unique PID
- Check process limits for user
- Set child state to “created”
- Copy data from parent PCB slot to child
- Increment counts on current directory & open files
- Copy parent context in memory (or set copy on write)
- Set child state to “ready to run”
- Wait for the scheduler to run the process

fork()



→ Read → no copy }
Write → "Buri baat"

(Create a) "Copy on write"

CoW

Fork Example

```
#include <stdio.h>

main(int argc, char **argv) {
    int pid;

    switch (pid=fork()) {
    case 0:    printf("I'm the child\n");
               break;
    default:
               printf("I'm the parent of %d\n", pid);
               break;
    case -1:
               perror("fork");
    }
}
```

```
int main( ... ) {
    int pid;
```


```
    pid = fork();
```

```
    if (pid == 0)
        printf("Child");
```

```
    else if (pid == -1)
        printf("Error");
```

```
    else
        printf("Parent of %d", pid);
```

```
}
```



$P_1 \rightarrow$ Child

$P_1 \rightarrow$ Parent of _____


Parent of _____
Child.

```
int main() {
```

```
    fork()
```

```
    fork()
```

```
    fork()
```



main()

int pid = 0;

pid = fork();

if (pid == 0)



fork() ←

fork()

else

fork() ↑

}



Running other programs

execve: replace the current process image with a new one

– See also *execl*, *execle*, *execlp*, *execvp*, *execvP*

- New program inherits:

- Processes group ID
- Open files
- Access groups
- Working directory
- Root directory
- Resource usages & limits
- Timers
- File mode mask
- Signal mask

Exec Example

```
#include <unistd.h>

main(int argc, char **argv) {
    char *av[] = { "ls", "-al", "/", 0 };

    execvp("ls", av);
}
```

Fork & exec combined

- UNIX creates processes via *fork* followed by *exec*
- Windows approach
 - *CreateProcess* system call to create a new child process
 - Specify the executable file and parameters
 - Identify startup properties (windows size, input/output handles)
 - Specify directory, environment, and whether open files are inherited

Exiting a process

exit system call

```
#include <stdlib.h>

main(int argc, char **argv) {
    exit(0);
}
```

exit: what happens?

- Ignore all signals
- If the process is associated with a controlling terminal
 - Send a hang-up signal to all members of the process group
 - reset process group for all members to 0
- close all open files
- release current directory
- release current changed root, if any
- free memory associated with the process
- write an accounting record (if accounting)
- make the process state zombie
- assign the parent process ID of any children to be 1 (init)
- send a “death of child” signal to parent process (SIGCHLD)
- context switch (we have to!)

Wait for a child process to die

wait system call

- Suspend execution until a child process exits
- *wait* returns the exit status of that child.

```
int pid, my_pid, status;

switch (my_pid=fork()) {
case 0:      /* do child stuff */ break;
case -1:     /* do error stuff */ break;

default:     /* wait for child to exit */
    while (pid=wait(&status))
        if (pid==my_pid)
            printf("got exit of %d\n", WEXITSTATUS(status));
            break;
}
```

```
int main () {
```

```
    int pid ;
```

```
    pid = fork() ;
```

```
    if ( pid == 0 ) {  
        |      "Child running"  
    } else {
```

```
        |      "Parent running"
```



```
        |      wait (NULL) ;
```

```
        |      "Child terminated"
```


```
    }
```

```
    "Done"
```

```
}
```



Parent & child processes



Signals


- Inform processes of asynchronous events
 - Processes may specify signal handlers
- Processes can poke each other (if they are owned by the same user)
- Sending a signal:
 - *kill* (*int pid, int signal_number*)
- Detecting a signal:
 - *signal* (*signal_number, function*)

<signal.h>

```
int main() {
```

```
    signal(SIGINT, handler);
```

```
    while (1);
```



```
}
```

```
void handler(int sig) {
```

```
    "Signal", sig ;
```

```
}
```

The End