

GETTING INFO ABOUT PROCESSES...


/proc
↳ "virtual filesystem"

cat /proc/cpuinfo
meminfo
interrupts

ls -l /proc/1

cat /proc/1/maps
status



/* status



x ————— x

"Light weight processes"

"Thread of execution"





Operating Systems Design

5. Threads

Paul Krzyzanowski
pxk@cs.rutgers.edu

Thread of execution

Single sequence of instructions

- Pointed to by the program counter (PC)
- Executed by the processor

Conventional programming model & OS structure:

- Single threaded
- One process = one thread

Multi-threaded model

A thread is a subset of a process:

- A process contains one or more ~~kernel~~ threads


Share memory and open files

- **BUT:**
separate program counter, registers, and stack
- Shared memory includes the heap and global/static data
- **No memory protection among the threads**

Preemptive multitasking:

- Operating system preempts & schedules threads

concurrent : actually in parallel
simultaneous : illusion of concurrency



Sharing

Threads share:


- ✓ Text segment (instructions)
- ✓ Data segment (static and global data)
- ✓ BSS segment (uninitialized data)
- ✓ Open file descriptors
- ✓ Signals
 - Current working directory
 - User and group IDs

Threads do not share:

- Thread ID
- Saved registers, stack pointer, instruction pointer
- Stack (local variables, temporary variables, return addresses)
- Signal mask
- Priority (scheduling information)



Android



Why is this good?

Threads are more efficient


- Much less overhead to create: no need to create new copy of memory space, file descriptors, etc.

Sharing memory is easy (automatic)

- No need to figure out inter-process communication mechanisms

Take advantage of multiple CPUs – just like processes

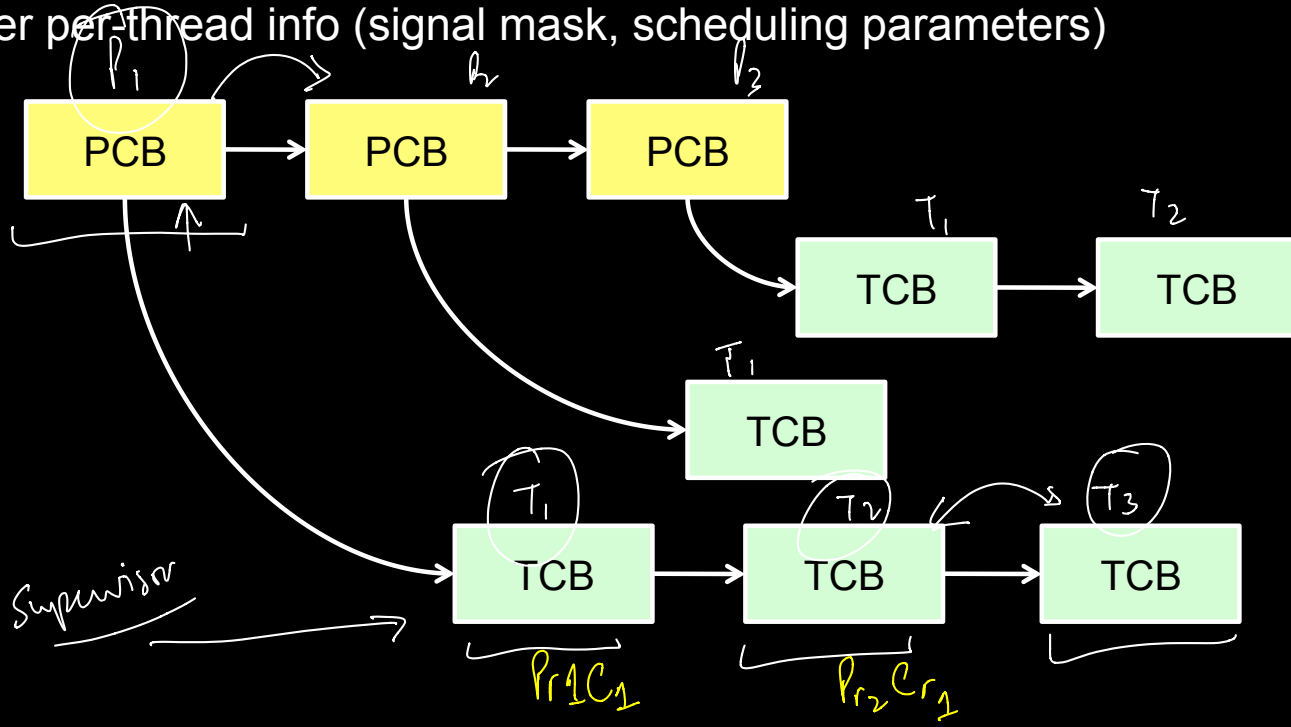
- Program scales with increasing # of CPUs
- Take advantage of multiple cores



Implementation

Process info (Process Control Block) contains one or more **Thread Control Blocks (TCB)**:

- Thread ID
- Saved registers
- Other per-thread info (signal mask, scheduling parameters)



Scheduling

Supervisor

¹¹
A threaded-aware operating system¹¹ scheduler schedules threads, not processes

- A process is just a container for one or more threads



Scheduler has to realize:


- Context switch among threads of different processes is more expensive:

- Flush cache memory (or have memory with process tags)
- Flush virtual memory TLB (or have tagged TLB)
- Replace page table pointer in memory management unit

- Scheduling threads onto a different CPU is more expensive
 - The CPU's cache may have memory used by the thread cached

- CPU affinity





" Work flow "

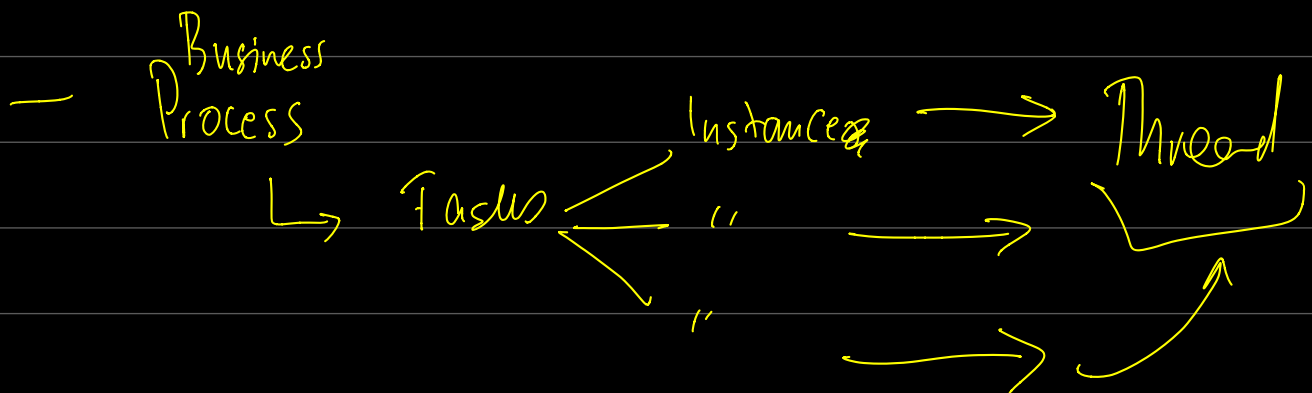
SAP

User requirements

Flows



Code

process



Multi-core architecture

CPU affinity



Programming patterns

Design Patterns

Single task thread

create - serve - die

- Do a specific job and then release the thread

Worker threads

create - serve - wait
→ specific

- Specific task for each worker thread
- Dispatch task to the thread that handles it

Thread pools

create - serve - wait




- ✈ - Create a pool of threads *a priori*
- Use an existing thread to perform a task; wait if no threads available
- Common model for servers

→ Programming Patterns.

HTTP 1.1
HTTP 2.0

- message passing.

When are threads created?
What do they do?



Kernel-level threads vs. User-level threads

Kernel-level

- Threads supported by operating system
- OS handles scheduling, creation, synchronization

User-level

- Library with code for creation, termination, scheduling
- Kernel sees one execution context: one process
- May or may not be preemptive

userspace

kernel space

User-level threads

Advantages

- ① – Low-cost: user level operations that do not require switching to the kernel
- Scheduling algorithms can be replaced easily & custom to app
- Greater portability

Disadvantages

- If a thread is blocked, all threads for the process are blocked
 - Every system call needs an asynchronous counterpart
- Cannot take advantage of multiprocessing


You can have both

User-level thread library on top of multiple kernel threads

$1:1$ – pure kernel threads only
(1 user thread = 1 kernel thread)

$N:1$ – pure user threads only
(N user threads on 1 kernel thread/process)

$N:M$ – hybrid threading
(N user threads on M kernel threads)



pthread: POSIX Threads

- POSIX.1c, Threads extensions (IEEE Std 1003.1c-1995)
- Defines API for managing threads
- Linux: native POSIX Thread Library (as of 2.6 kernel)
- Also on Solaris, Mac OS X, NetBSD, FreeBSD
- API library on top of Win32

"pthread.h not found"

Using POSIX Threads

Create a thread *struct*

`pthread_t t;`

`pthread_create(&t, NULL, func, arg)`

function to execute

job/task/service

- Create new thread *t*
- Start executing function *func(arg)*

Join two threads:

`void *ret_val;`

`pthread_join(t, &ret_val);`

returned value

- Wait for thread *t* to terminate (via *return* or *pthread_exit*)

No parent/child relationship!

- Any one thread may wait (join) on another thread

Linux *clone()* system call

- Clone a process, like *fork*, but:
 - Specify function that the child will run (with argument)
 - Child terminates when the function returns
 - Specify location of the stack for the child
 - Specify what's shared:
 - Share memory (otherwise memory writes use new memory)
 - Share open file descriptor table
 - Share the same parent
 - Share root directory, current directory, and permissions mask
 - Share namespace (mount points creating a directory hierarchy)
 - Share signals
 - *And more...*
- Used by pthreads

Threading in hardware


(Not included)

- Hyper-Threading (HT) vs. Multi-core vs. Multi-processor

- One core = One CPU

- Hyper-Threading

- One physical core *appears* to have multiple processors
 - Looks like multiple CPUs to the OS
- Multiple threads run but compete for execution unit
- Events in the pipeline switch between the streams
- Threads do not have to belong to the same process
 - But the processors share the same cache
 - Performance can degrade if two threads compete for the cache
- Works well with instruction streams that have large memory latencies



Amdahl's Law

100,000,000
↓
100ns

500,000
↓
50ns

500,000
↓
50ns

10000...
↓
ns

The End