

Operating Systems Design

7. Process Scheduling

Paul Krzyzanowski
pxk@cs.rutgers.edu

Process Behavior

Most processes exhibit:


- ~~Large~~ # of short CPU bursts between I/O requests
- ~~Small~~ # of long CPU bursts between I/O requests

→ Interactive

→ computational


→ I/O bound processes

↓
CPU-bound processes




Process Behavior

Interactive process: mostly short CPU bursts



Compute process: mostly long CPU bursts

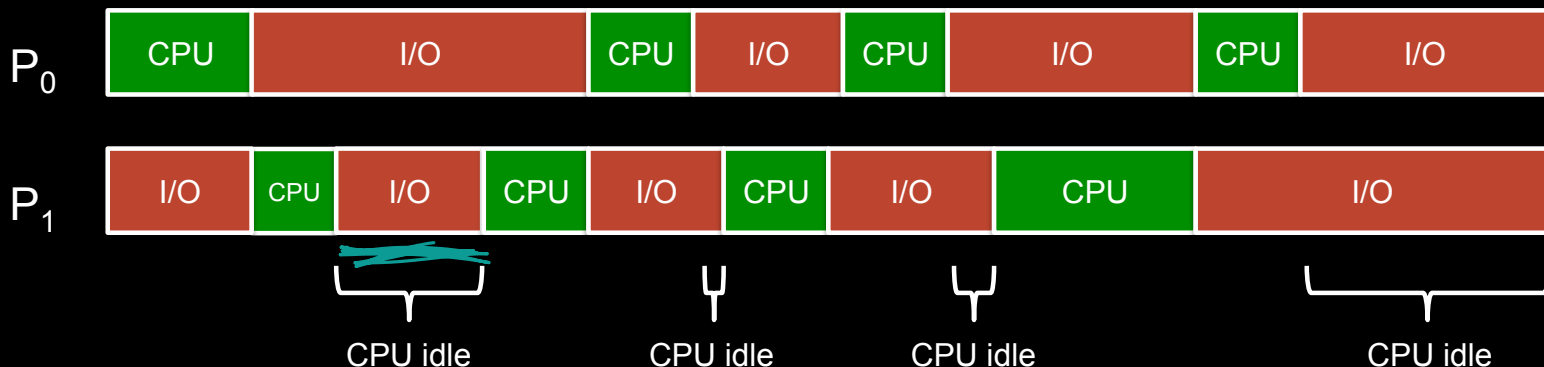


Process Scheduling

Goal:

- Maximize use of CPU & improve throughput
- Let another process run when the current one is waiting on I/O

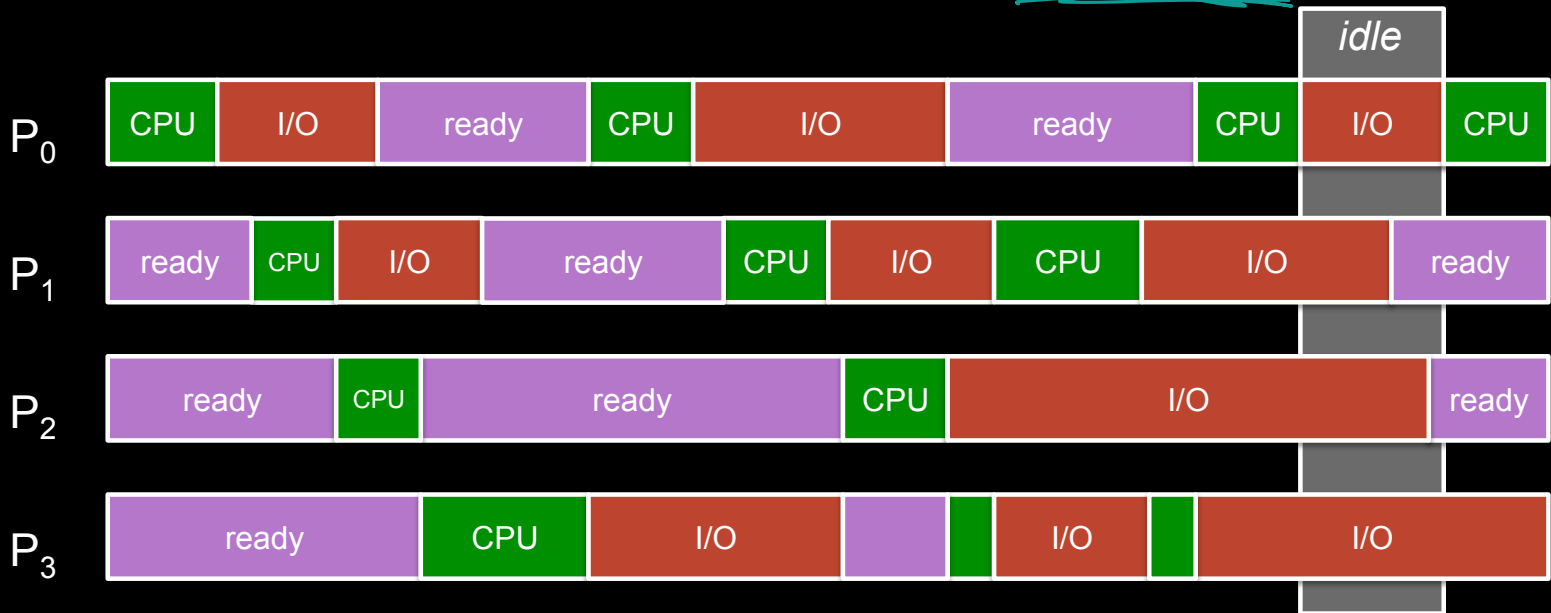
of processes completing their jobs.



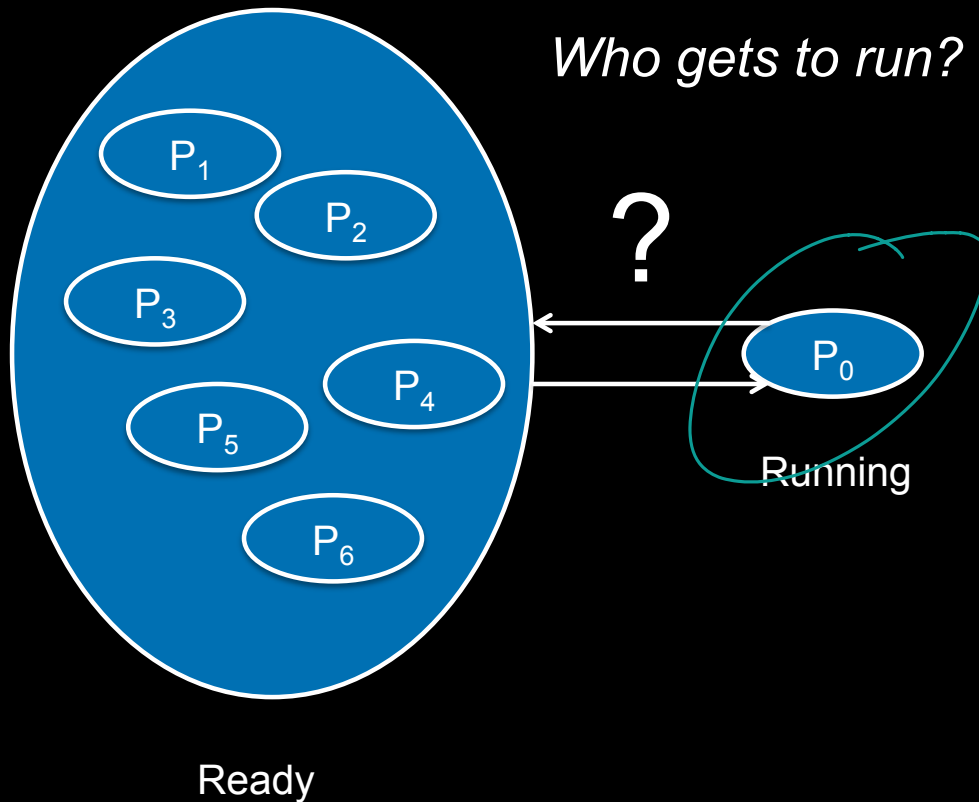
Process Scheduling

Reality:

- Some processes will use long stretches of CPU time
 - Preempt them and let another process run
- More processes may want the CPU: keep them in the *ready* list
- Perhaps all processes are waiting on I/O: *nothing to run!*



Process Scheduler



Switching processes

- Scheduling algorithm:

- Policy: Makes the decision of who gets to run

- Dispatcher:

- Mechanism to do the context switch

} context switch

When does the scheduler make decisions?

Four events affect the decision:

1. Current process goes from *running* to *waiting* state ←
2. Current process terminates ←
3. Interrupt causes the scheduler to move a process from *running* to *ready*:
scheduler decides it's time for someone else to run → *preemption*
4. Current process goes from *waiting* to *ready* ← *I/O completed*.
I/O (including blocking events, such as semaphores) is complete

- Preemptive scheduler vs.

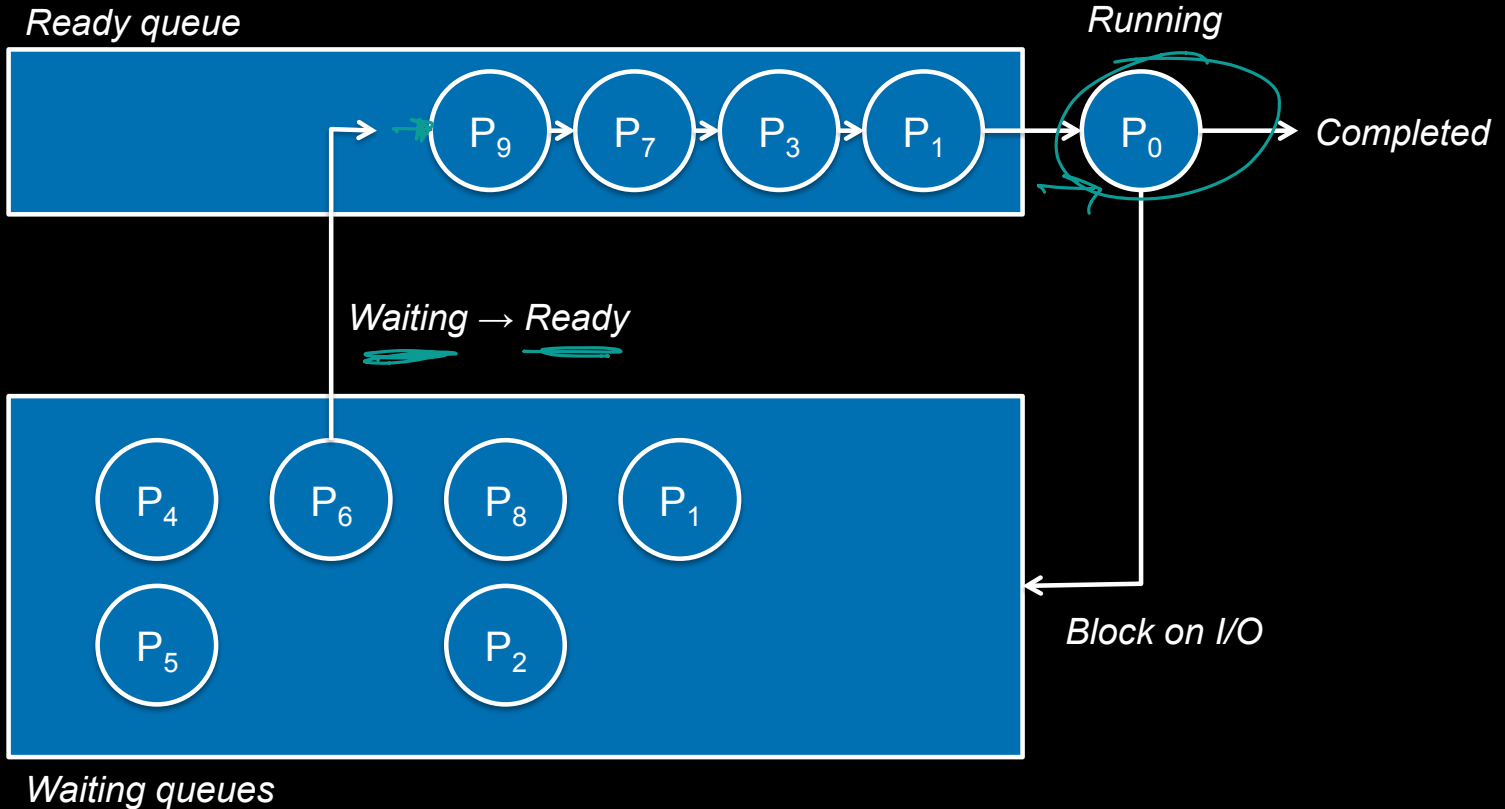
- Cooperative (non-preemptive) scheduler
 - CPU cannot be taken away

- Run-to-completion scheduler (old batch systems)

Scheduling algorithm goals

- Be fair (to processes? To users?)
- Be efficient: Keep CPU busy ... and don't spend a lot of time deciding!
- Maximize throughput: minimize time users must wait
- Minimize response time *Process start → first response*
- Be predictable: jobs should take about the same time to run when run multiple times
- Minimize overhead
- Maximize resource use: try to keep devices busy! ↓
- Avoid starvation
- Enforce priorities
- Degrade gracefully Graceful degradation

First-Come, First-Served (FCFS)

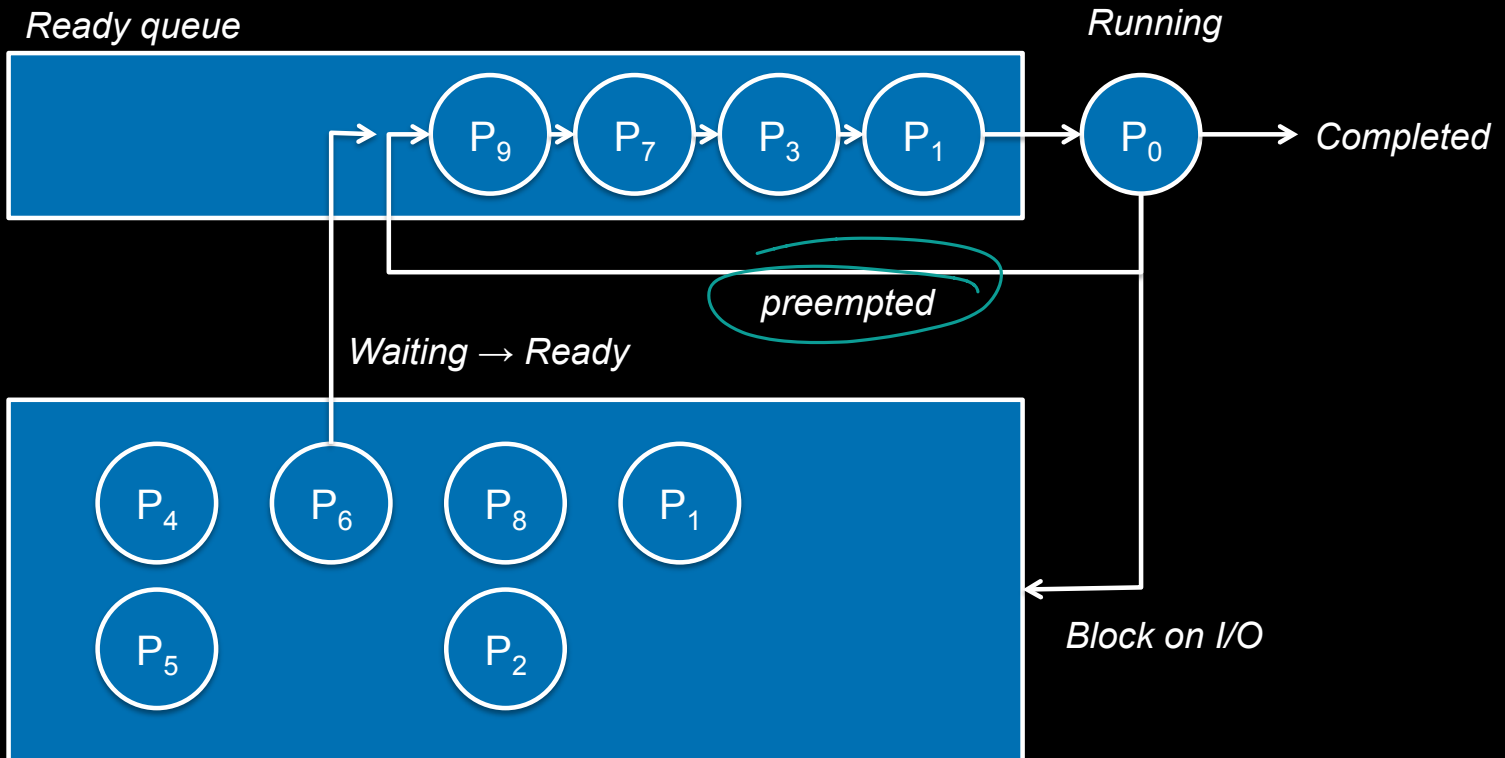


First-Come, First-Served (FCFS)

- Non-preemptive ✓
- A process with a long CPU burst will hold up other processes
 - I/O bound jobs may have completed I/O and are ready to run: poor device utilization
 - Poor average response time

Round-Robin Scheduling

Preemptive: Process can not run for longer than a quantum (time slice)



Waiting queues

Round-Robin Scheduling

- Performance depends on the time slice
 - Long time slice makes this similar to FCFS
 - Short time slice increases overhead % of context switching
- Advantages
 - Every process gets an equal share of the CPU ←
 - Easy to implement
 - ✓ Easy to compute ~~average response time~~: $f(\# \text{ processes on list})$
- Disadvantage
 - Giving every process an equal share isn't necessarily good
 - Highly interactive processes will get scheduled the same as CPU-bound processes

Shortest Remaining Time First Scheduling

- Sort jobs by anticipated CPU burst time
- Schedule shortest ones first
- Optimize average response time

Burst time	2	2	10	3	8
Process	E	D	C	B	A
Total run time	25	23	21	11	8
	last				first

Total time = 25

Mean time = 17.6

4 8
B 11

Burst time	10	8	3	2	2
Process	C	A	B	D	E
Total run time	25	15	7	4	2

Total time = 25

Mean time = 10.6

Mean completion time for a process falls by almost 40%!

Shortest Remaining Time First Scheduling

- Biggest problem: we're optimizing with data we don't have!

- All we can do is estimate

- Exponential average:

$$e_{n+1} = \alpha t_n + (1 - \alpha)e_n$$

α is a weight factor to balance the weight of the last burst period vs. historic periods ($0 \leq \alpha \leq 1$)

If $\alpha = 0$: $e_{n+1} = e_n$ (recent history has no effect)

If $\alpha = 1$: $e_{n+1} = \alpha t_n$ (use only the last burst time)

- Algorithm can be preemptive or non-preemptive
- Preemptive version is:

Shortest remaining time first scheduling (vs. SJF)

Shortest Job first

Shortest Remaining Time First Scheduling

- **Advantage**

- Short-burst jobs run fast

- **Disadvantages**

- ~~✗~~ Long burst (CPU intensive) jobs get a long mean waiting time
- ✓ – Rely on ability to estimate CPU burst length

Priority Scheduling

Round Robin assumes all processes are equally important

nice

- Not true
 - Interactive jobs need high priority for good response
 - Long non-interactive jobs can worse treatment (get the CPU less frequently): *this goal led us to SRTF*
 - Users may have different status (e.g., administrator)
- **Priority scheduling** algorithm:
 - Each process has a priority number assigned to it
 - Pick the process with the highest priority
 - Processes with the same priority are scheduled round-robin

Priority Scheduling

- Priority assignments:
 - **Internal**: time limits, memory requirements, I/O:CPU ratio, ...
 - **External**: assigned by administrators
- Static & dynamic priorities
 - **Static priority**: priority never changes
 - **Dynamic priority**: scheduler changes the priority during execution
 - Increase priority if it's I/O bound for better interactive performance or to increase device utilization
 - Decrease a priority to let lower-priority processes run
 - Example: use priorities to drive SJF/SRTF scheduling

Priority Scheduling: dealing with starvation

- **Starvation**
 - Process is blocked indefinitely
 - Steady stream of higher-priority processes keeps it from being scheduled
- Dealing with starvation: **Process aging**
 - Gradually increase the priority of a process so that eventually its priority will be high enough so it will be scheduled to run
 - Then bring it down again

Multilevel Queues


- **Priority classes**

- Examples: System processes, interactive processes, slow interactive processes, background non-interactive processes
- Each priority class gets its own queue
- Processes are permanently assigned to a specific queue

- **Goals**

- Priority scheduler with queues per priority level
- Each queue may have a different scheduling algorithm
- Quantum is increased at each lower priority level
 - Lower-priority processes tend to be compute bound

Multilevel Queues



Multilevel Feedback Queues

- **Advantage**

- Good for separating processes based on CPU burst needs
- Let I/O bound processes run often
- Give CPU-bound processes longer chunks of CPU
- No need to estimate interactivity! (Estimates were often flawed)

- **Disadvantages**

- Priorities get controlled by the system.
A process is considered important because it uses a lot of I/O
- Processes whose behavior changes may be poorly scheduled
- System can be gamed by scheduling bogus I/O


Symmetric multiprocessor scheduling

- **Processor affinity**


- Try to reschedule a process onto the same CPU
- Cached memory may be present on the CPU's cache

- **Types of affinity**

- **Hard** : force a process to stay on the same CPU
- **Soft affinity**: best effort, but the process may be rescheduled on a different CPU
 - **Load balancing**: ensure that CPUs are busy
 - It's better to run a job on another CPU than wait
 - If the run queue for a CPU is empty, get a job from another CPU's run queue: **pull migration**
 - Check load periodically: if not balanced, move jobs. **Push migration**



Scheduler Examples



Solaris Scheduler

- Priority-based scheduler: 170 priorities (0-169)
 - High priority → short quantum
- Six scheduling classes
 - Each class has priorities and scheduling algorithms

1. Time sharing (0-59)

Default class. Dynamic priorities via a multilevel feedback queue *DEFAULT*

2. Interactive (0-59)

Like TS but higher priority for in-focus windows in GUI

3. Real-time (100-159)

Fixed priority, fixed time quantum; high priority values

4. System (60-99)

Used to schedule kernel threads: run until they block or complete

5. Fair share (0-59)

Processes scheduled on % of CPU

6. Fixed priority (0-59)

Fixed priority

Highest priority (160-169): interrupt-handling threads

Solaris Scheduler

- Default class: **time sharing**
 - Multilevel feedback queue
 - Small time slice for high priority queue
 - Long time slice for low priority queue
- **Interactive class**: similar but gives windowing apps higher priority
- Highest priority: threads in the **real-time class**
- **System class**: runs kernel threads (scheduler & paging)
 - Not preempted
- **Fair share**: set of processes get a “CPU share”
- **Fixed priority**: like time-sharing but never adjusted

Windows Scheduler

- Two classes:
 - Variable class: priorities 1-15
 - Real-time class: priorities 16-31
- Each priority level has a queue
 - Pick the highest priority thread that is ready to run
- Relative priority
 - Threads have relative levels within their class
 - When a quantum expires, the thread's priority is lowered but never below the base
 - When a thread wakes from wait, the priority is increased
 - Higher increase if waiting for keyboard input
 - Priority is increased for foreground window processes

Linux Schedulers

- Linux 1.2: Round Robin scheduler (fast & simple)
- Linux 2.2: Scheduling classes
 - Classes: Real-time, non-real-time, non-preemptible
 - Support for symmetric multiprocessing
- Linux 2.4: O(N) scheduler
 - Iterates over every task at each scheduling event
 - If a time slice was not fully used, $\frac{1}{2}$ of the remaining slice was added to the new time slice for the process.
 - “goodness” metric decided who goes next
 - One queue (in a mutex): no processor affinity

Linux 2.6 $O(1)$ scheduler goals

Addressed three problems

- Scalability: $O(1)$ instead of $O(n)$ to not suffer under load
- Support processor affinity
- Support preemption

Linux 2.6 O(1) scheduler

- One runqueue per CPU: 140 priority lists serviced round robin
 - Two priority ranges: 0-99 for real-time; 100-140 for others
 - *High priority processes get a longer quantum!*
 - If a process uses its time slice, it will not get executed until all other processes exhaust their quanta
- runqueue data structure:
 - Two arrays sorted by priority value:
 - **Active**: all tasks with time remaining in their slices
 - **Expired**: all tasks that used up their time slice
 - Scheduler chooses the highest priority task from the active queue
 - When the active queue is empty, the expired queue becomes active

Linux 2.6 O(1) scheduler

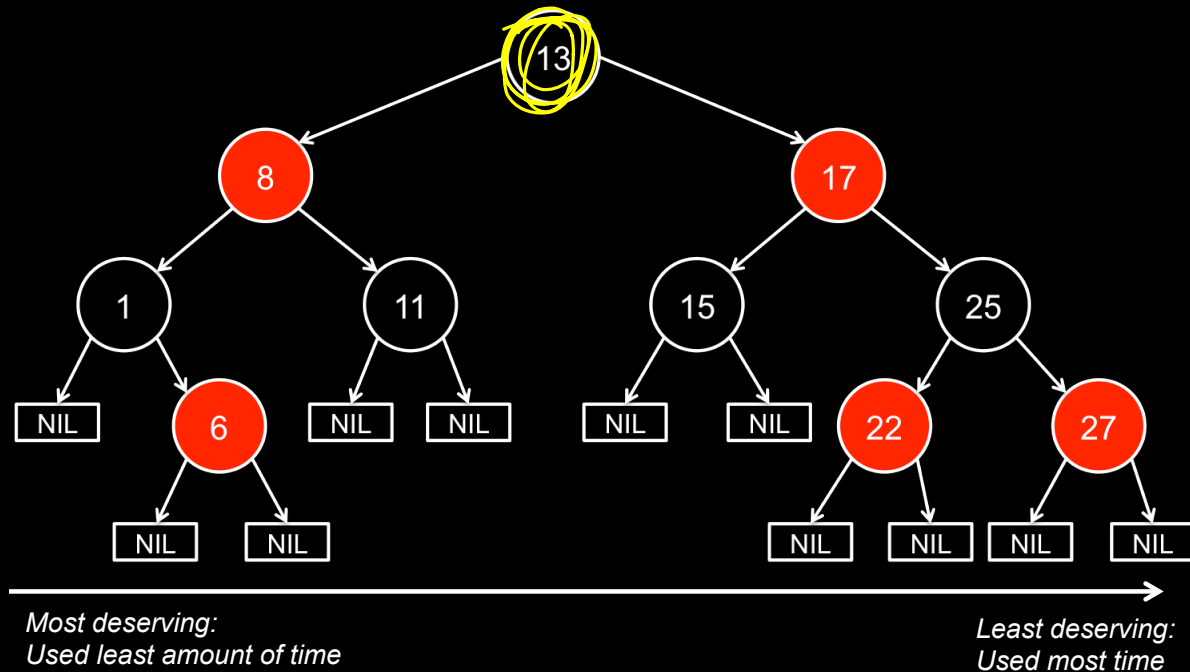
- Real-time tasks: static priorities
- Non real-time tasks: dynamic priorities
 - I/O-bound processes get priority increased by up to 5 levels
 - CPU-bound processes get priority decreased up to 5 levels
 - Interactivity determined by %sleep : %compute time ratio
- SMP load balancing
 - Every 200ms, check if CPU loads are unbalanced
 - If so, move tasks from a loaded CPU to a less-loaded one
 - If a CPU's runqueue is empty, move from the other runqueue
- Downside of O(1) scheduler
 - A lot of code with complex heuristics

Linux Completely Fair Scheduler

- Latest scheduler (introduced in 2.6.23)
- Goal: give a “fair” amount of CPU time to tasks
- Keep track of time given to a task (“virtual runtime”)
 - Also use “**sleeper fairness**”: tasks get a “fair” share of the CPU even if they sleep from time to time
- Priorities
 - Used as a decay factor for the time a task is permitted to execute
 - Allowable time decreases for low priority tasks

Linux Completely Fair Scheduler


- No run queues
- Time-sorted read-black tree instead of a run queue
 - Self-balancing binary tree: search, insert, & delete in $O(\log n)$



From: http://en.wikipedia.org/wiki/File:Red-black_tree_example.svg

Linux Completely Fair Scheduler

- Goal: give a “fair” amount of CPU time to tasks
- Keep track of time given to a task (“virtual runtime”)
 - Also “sleeper fairness”: tasks that are waiting receive a fair share of the CPU when they are ready
- Time-sorted read-black tree instead of a run queue

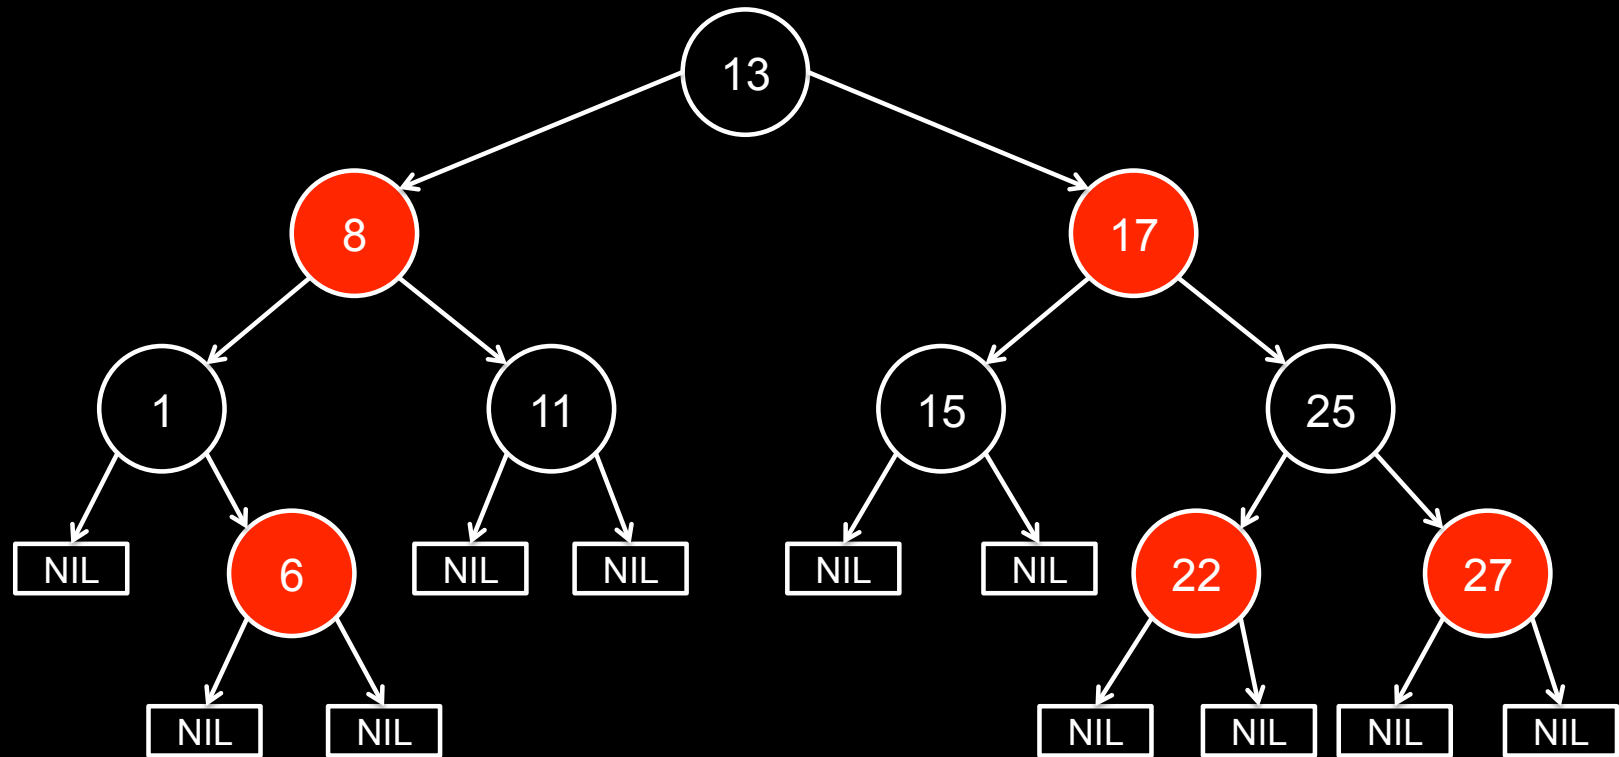


*Red-black tree:
self-balancing binary tree
 $O(\log n)$*

*Most deserving:
Used least amount of time*

*Least deserving:
Used most time*

From: http://en.wikipedia.org/wiki/File:Red-black_tree_example.svg



CFS: picking a process

- Scheduling decision:
 - Pick the leftmost task
- When a process is done:
 - Add execution time to the per-task run time count
 - Insert the task back in the queue
- Heuristic: *decay factors*
 - Determine how long a task can execute
 - Higher priority tasks have lower factors of decay.
 - Avoids having run queues per priority level

Group Scheduling

- Default operation: be fair to each task
- Assign one virtual runtime to a group of processes
 - Per user scheduling
 - cgroup pseudo file system interface for configuring groups
 - E.g., a user with 5 processes can get the same % of CPU as a user with 50 processes
- Default task group: `init_task_group`
- Improve interactive performance
 - A task calls `__proc_set_tty` to move to a tty task group
- `/proc/sys/kernel/sched_granularity_ns`
 - Tunable parameter to tune the scheduler between desktop (highly interactive) and server loads

More on the Linux scheduler

- Modular scheduler core: Scheduling classes
 - Scheduling class defines common set of functions that define the behavior of that scheduler
 - Add a task, remove a task, choose the next task
 - Each task belongs to a scheduling class
 - sched_fair.c
 - implements the CFS scheduler
 - sched_rt.c
 - implements a priority-based round-robin real-time scheduler
- Scheduling domains
 - Group one or more processors hierarchically
 - One or more processors can share scheduling policies

The End