



Operating Systems Design

8. Memory Management

Paul Krzyzanowski
pxk@cs.rutgers.edu

CPU Memory Access


CPU reads instructions and reads/write data from/to memory



Functional interface:

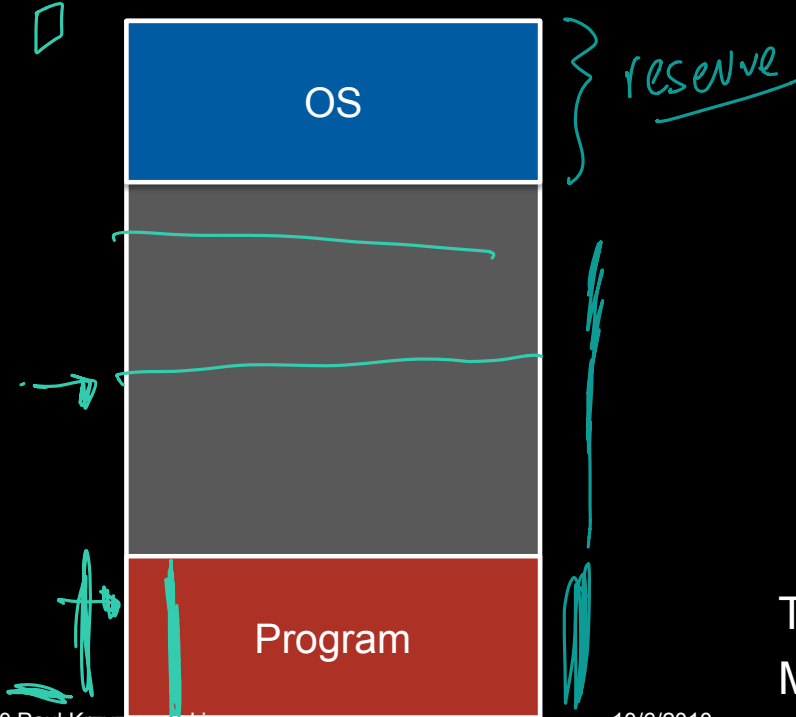
`value = read(address)` ←

`write(address, value)` ←



Monoprogramming


- Run one program at a time
- Share memory between the program and the OS



This was the model in old MS-DOS systems

Multiprogramming

- Keep more than one process in memory
- More processes in memory improves CPU utilization



- If a process spends 20% of its time computing, then would switching among 5 processes give us 100% CPU utilization?
- Not quite. For n processes, if $p = \% \text{ time a process is blocked on I/O}$ then:
$$\text{probability all are blocked} = p^n$$
- CPU is not idle for $(1-p^n)$ of the time
- 5 processes: 67% utilization

How do you access memory?

- **Absolute code**

if you know where the program gets loaded (any relocation is done at link time)

- **Position independent code**

all addresses are relative

- **Dynamically relocatable code**
relocated at load time


- Or ... use **logical addresses**

absolute code with addresses translated at runtime

.data
→


0x29ACD9 → RAM

Logical addressing



- **Memory management unit (MMU):**

Real-time, on-demand translation between
logical and *physical* addresses




CS register

DS register

mov [ESI], 25

mov [DS:ESI], 25

$$100 + 20 = 120$$

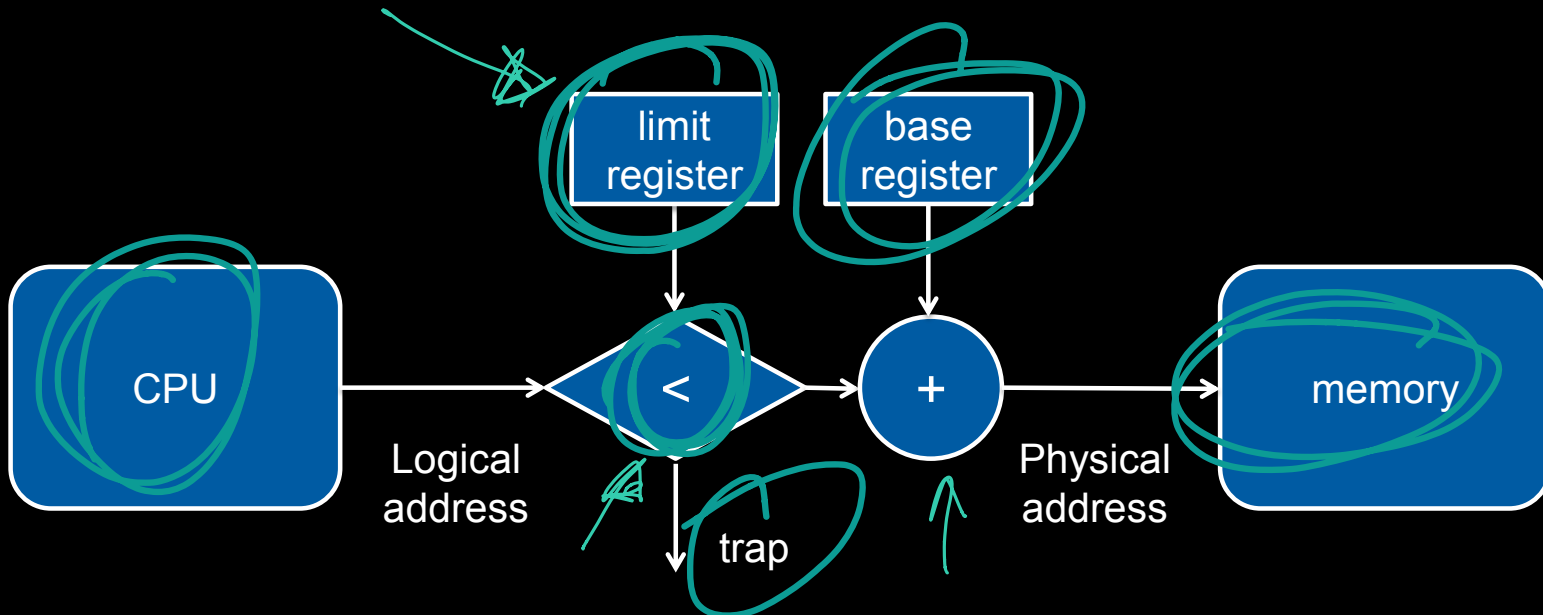



Relocatable addressing

15: 20 25
 $100 + 20 = 120$

- Base & limit

- $\text{Physical address} = \text{logical address} + \text{base register}$
- But first check that: $\text{logical address} < \text{limit}$

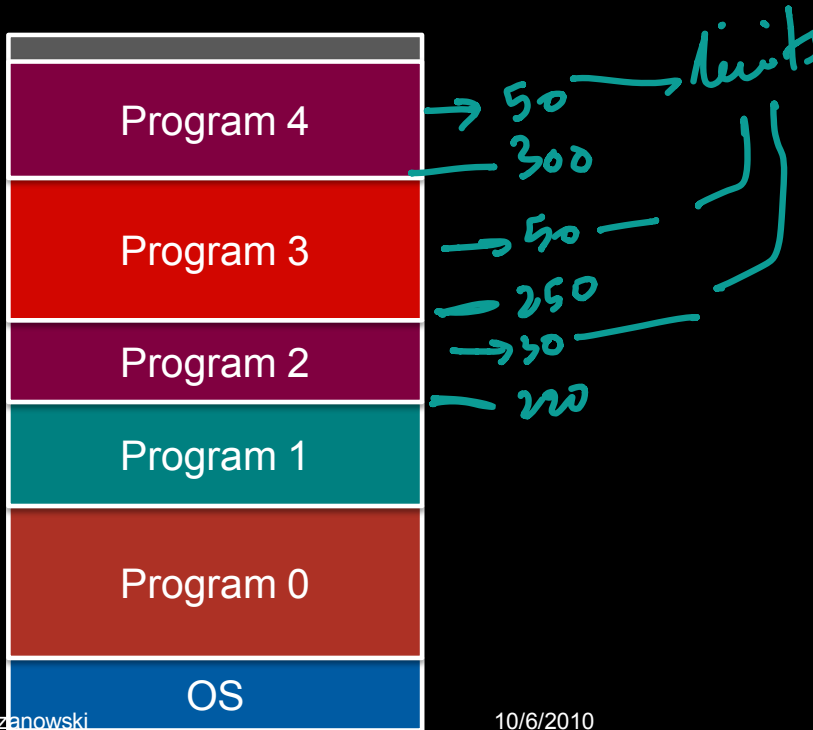


Multiple Fixed Partitions

- Divide memory into predefined partitions (segments)
 - Partitions don't have to be the same size
 - For example: a few big partitions and many small ones
- New process gets queued for a partition that can hold it
- Unused memory in a partition goes unused

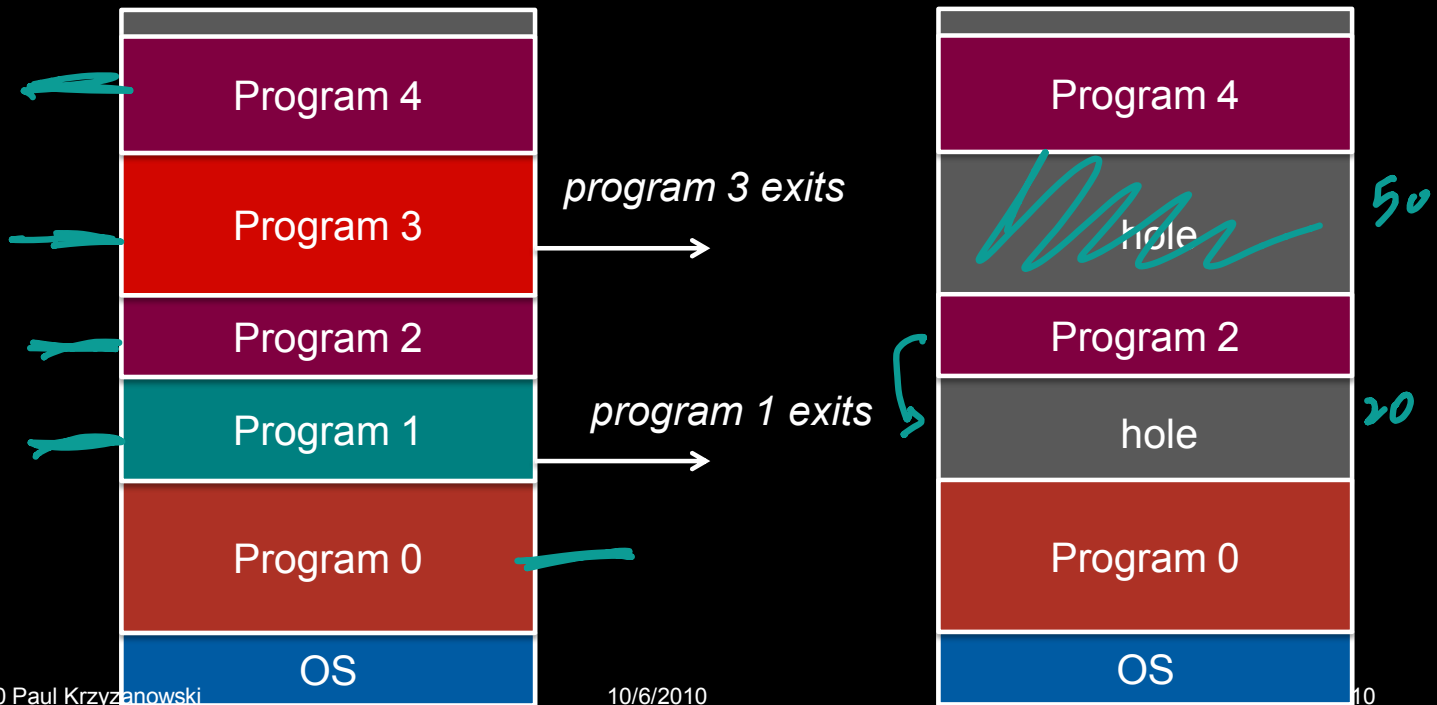
Variable partition multiprogramming

- Create partitions as needed
- New process gets queued
- OS tries to find a hole for it





Variable partition multiprogramming

- Create partitions as needed
- New process gets queued
- OS tries to find a hole for it

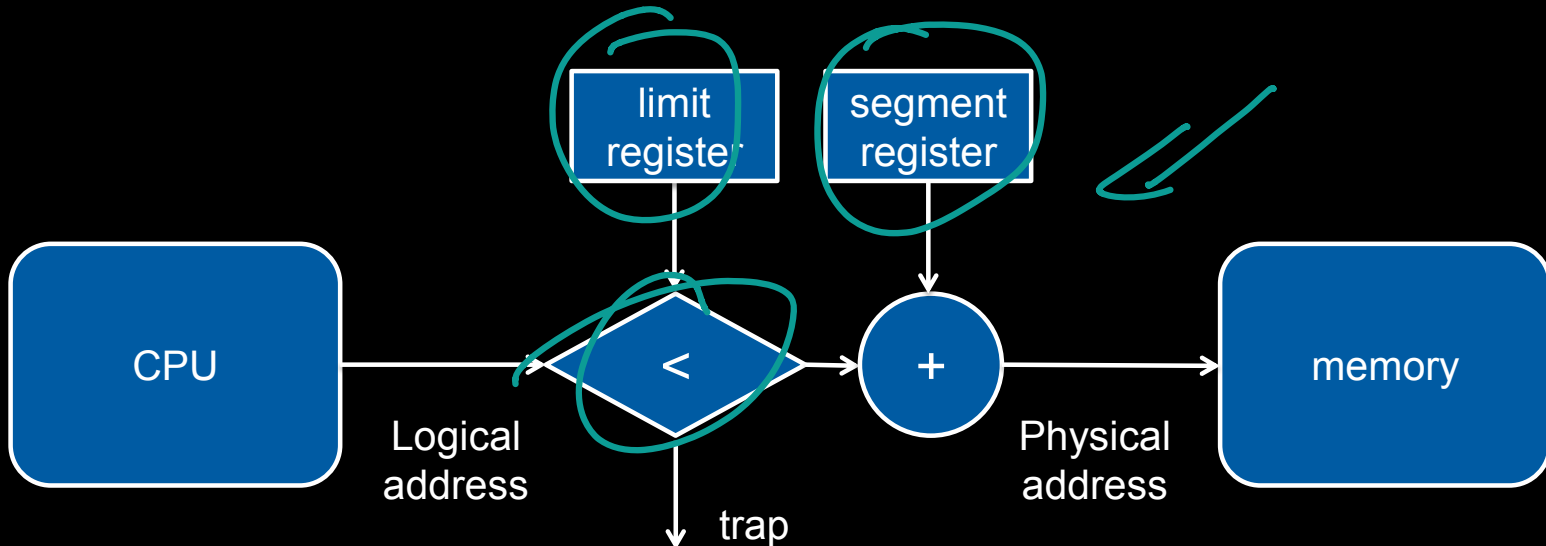


Variable partition multiprogramming

- What if a process needs more memory? 
 - Always allocate some extra memory just in case
 - Find a hole big enough to relocate the process
- Combining holes 
 - Memory compaction
 - Usually not done because of CPU time to move a lot of memory


Segmentation hardware

- Divide a process into segments and place each segment into a partition of memory
 - Code segment, data segment, stack segment, etc.




Allocation algorithms

- **First fit**: find the first hole that fits
- **Best fit**: find the hole that best fits the process
- **Worst fit**: find the largest available hole
 - *Why?* Maybe the remaining space will be big enough for another process. In practice, this algorithm does not work well.




Paging

- Memory management scheme
 - Physical space can be non-contiguous
 - No fragmentation problems
 - No need for compaction



Paging is implemented by the Memory Management Unit (MMU) in the processor



Paging


- Translation:

- Divide physical memory into fixed-size blocks: **page frames**
- A logical address is divided into blocks of the same size: **pages**
- All memory accesses are translated: **page** → **page frame**
- A page table maps pages to frames

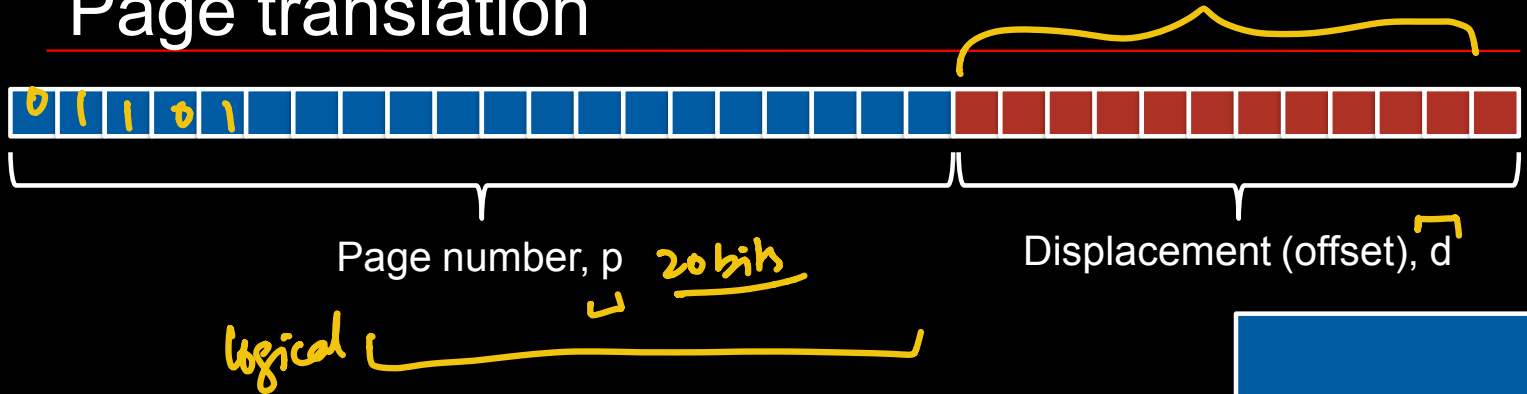
- Example:

- 32-bit address, 4 KB page size:
 - Top 20 bits identify the page number
 - Bottom 12 bits identify offset within the page/frame

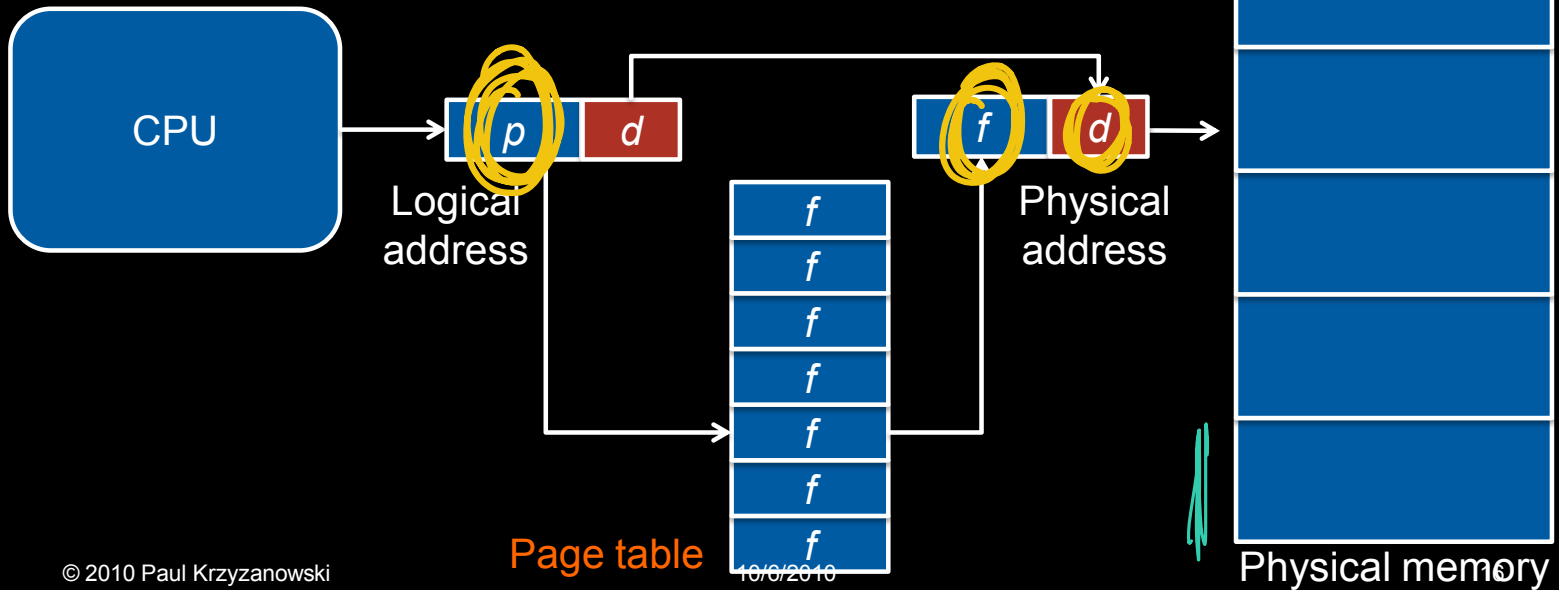
12-bit
= 4096



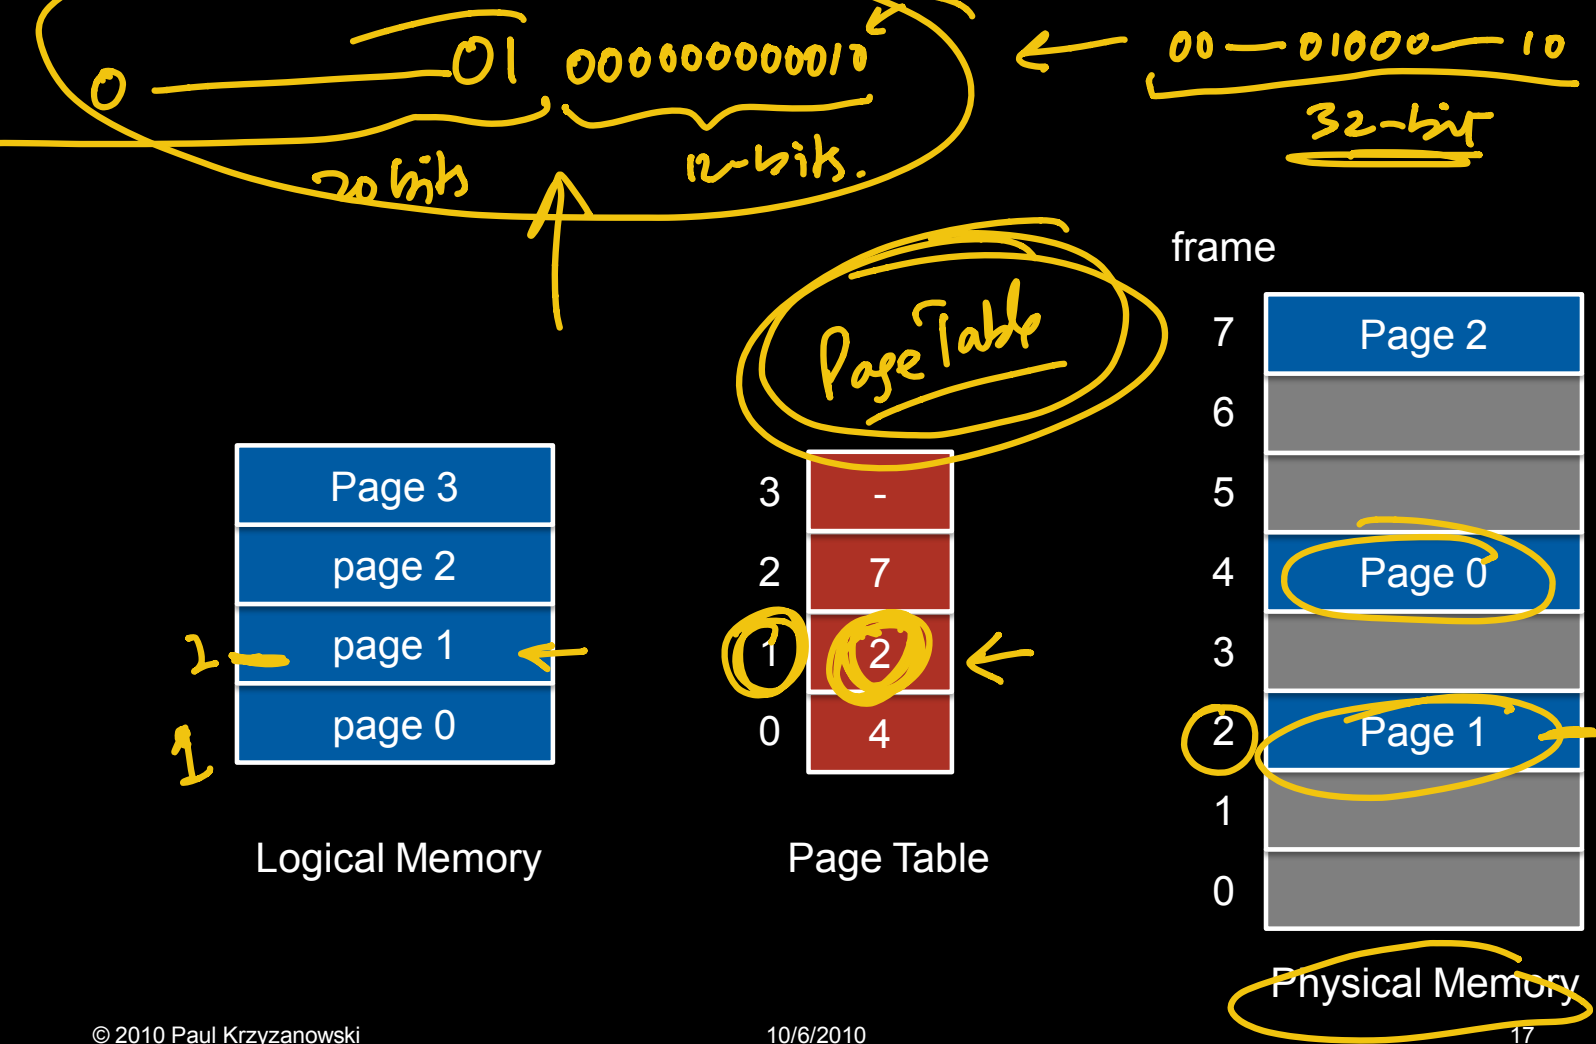
Page translation



$f = \text{page_table}[p]$





Logical vs. physical views of memory




Hardware Implementation

- Where do you keep the page table? *In memory*
- Each process gets its own virtual address space
 - Each process has its own page table
 - Change the page table by changing a *page table base register*
- Memory translation is now slow!
 - To read a byte of memory, we need to read the page table first
 - Each memory access is now 2x slower!




Hardware Implementation: TLB


- Cache frequently-accessed pages
 - Translation lookaside buffer (TLB)
 - Associative memory: key (page #) and value (frame #)
- TLB is on-chip & fast ... but small (64-1,024 entries)
- TLB miss: result not in the TLB
 - Need to do page table lookup in memory
- Hit ratio = % of lookups that come from the TLB
- Issue when we context switch




Process image




Process image



How does
malware
work?



- ① Writing ✓
- ② Executing ✓



Protection

- An MMU can enforce memory protection
- Page table stores protection bits per frame
 - Valid/invalid: is there a frame mapped to this page?
 - Read-only
 - No execute
 - Dirty

The End