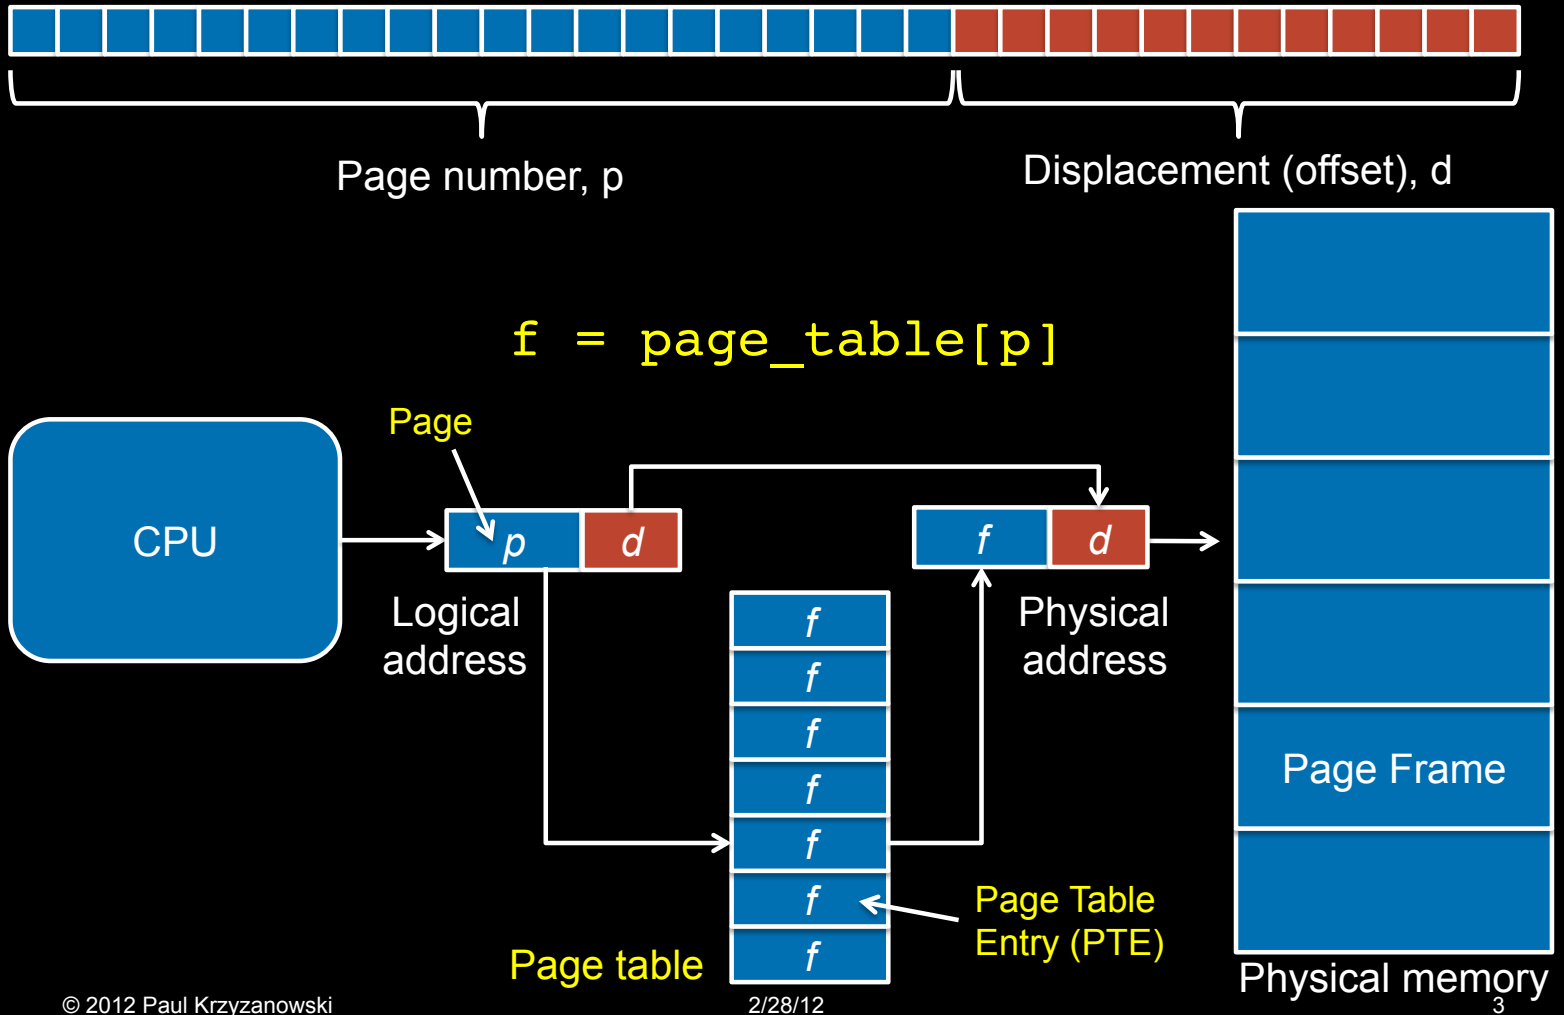


Operating Systems Design

10. Memory Management: Paging

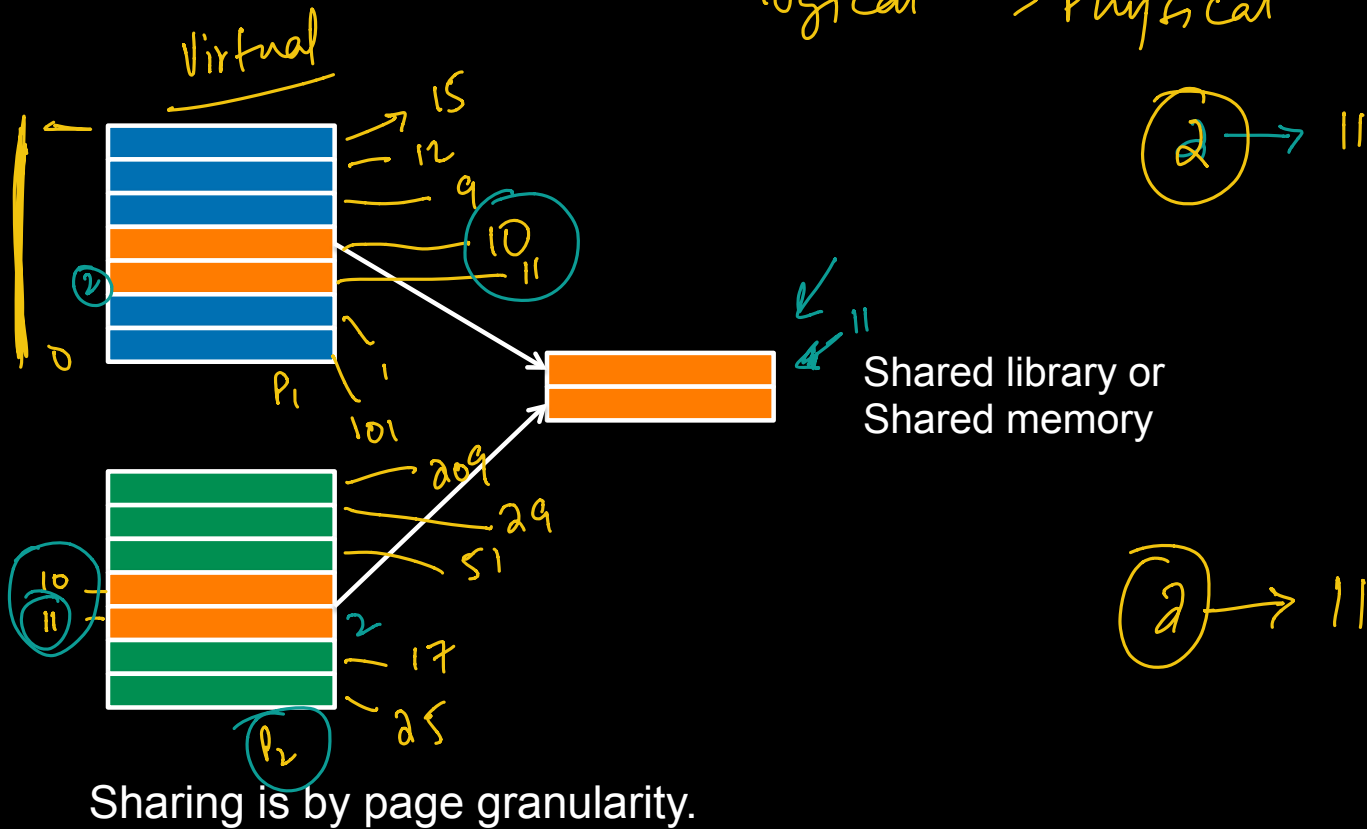
Paul Krzyzanowski
pxk@cs.rutgers.edu

Page translation

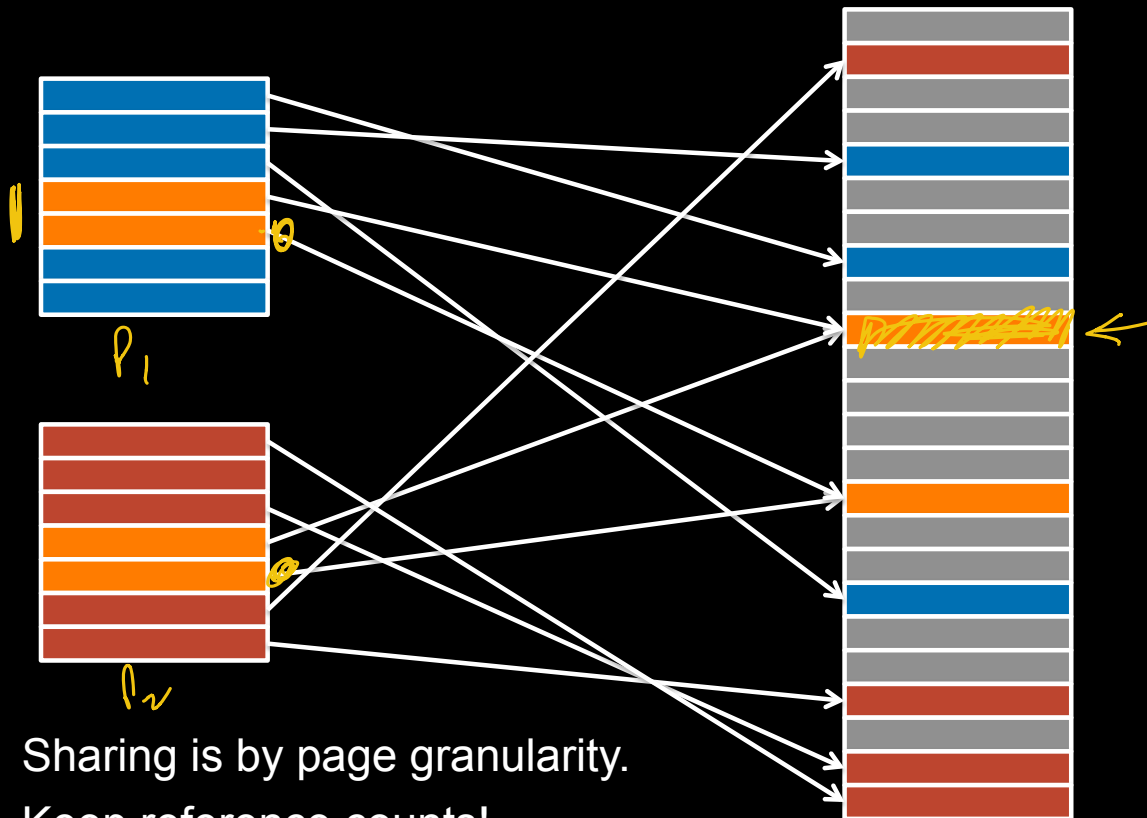


Virtual memory makes memory sharing easy

logical \rightarrow Physical



Virtual memory makes memory sharing easy



Sharing is by page granularity.
Keep reference counts!

dll

.so

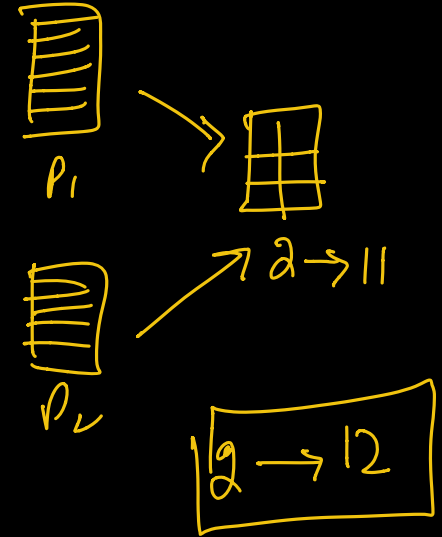
 .so

 .so

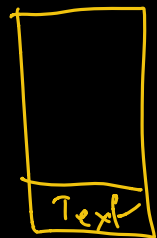
-lpthread

Copy on write

- Share until a page gets modified
- Example: fork()
 - Set all pages to read-only
 - Trap on write ←
 - If legitimate write
 - Allocate a new page and copy contents



Q: Do we keep the whole process image in RAM?



Demand Paging

Executing a program

- Allocate memory + stack and load the entire program from memory (including linked libraries)
- Then execute it

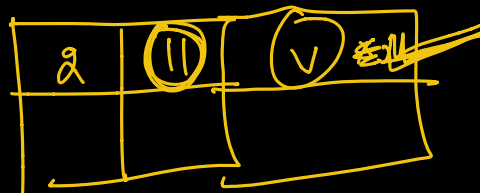
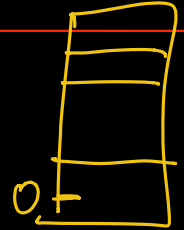
Executing a program

- Allocate memory + stack and load the entire program from memory (including linked libraries)
- Then execute it

We don't need to do this!

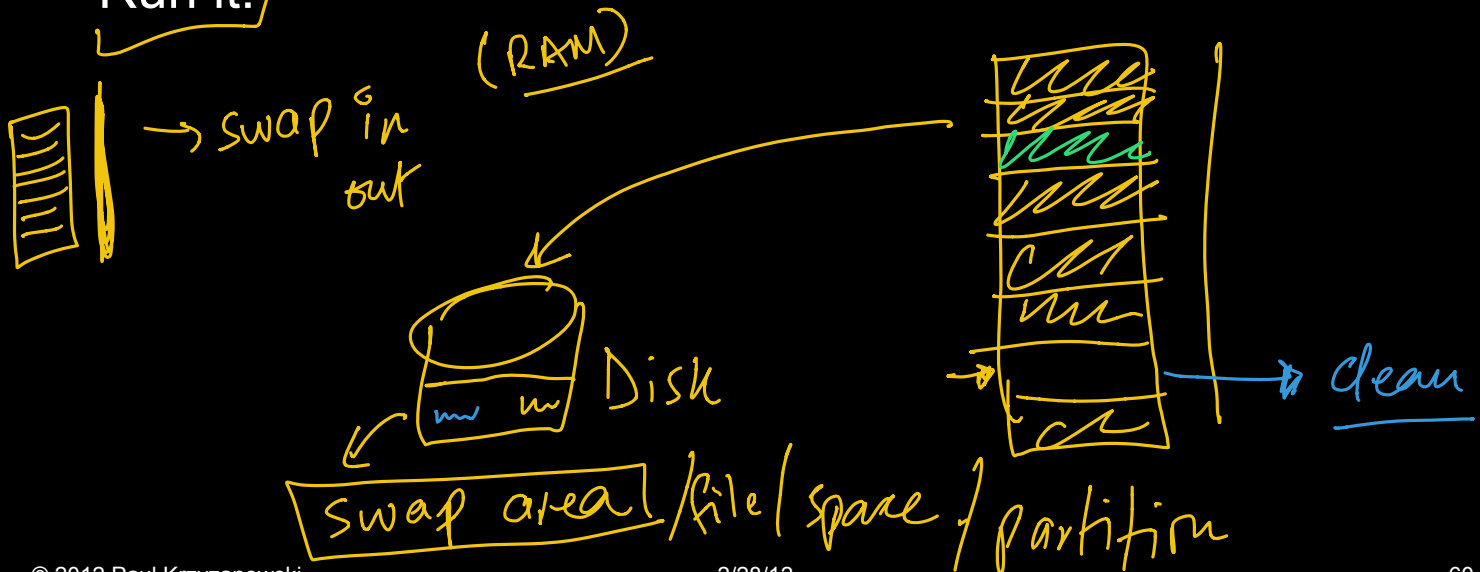
Demand Paging

- Load pages into memory only as needed ←
 - On first access
 - Pages that are never used never get loaded
- Use valid/invalid bit in page table entry
 - Valid: the page is in memory (“valid” mapping)
 - Invalid: ~~out of bounds access~~ or page is not in memory
 - ~~Have to check the process’ memory map in the PCB to find out~~
- Invalid memory access generates a **page fault**



Demand Paging: At Process Start

- Open executable file
- Set up memory map (stack & text/data/bss)
 - But don't load anything!
- Load first page & allocate initial stack page
- Run it!



Demand Paging: Page Fault Handling

- Soon the process will access an address without a valid page
 - OS gets a page fault from the MMU
- What happens?
 - Kernel searches a tree structure of memory allocations for the process to see if the faulting address is valid
 - If not valid, send a SEGV signal to the process
 - Is the type of access valid for the page?
 - Send a signal if not
 - We have a valid page but it's not in memory
 - Bring page in, mark as valid. ←
restart instruction

Page Replacement Policy

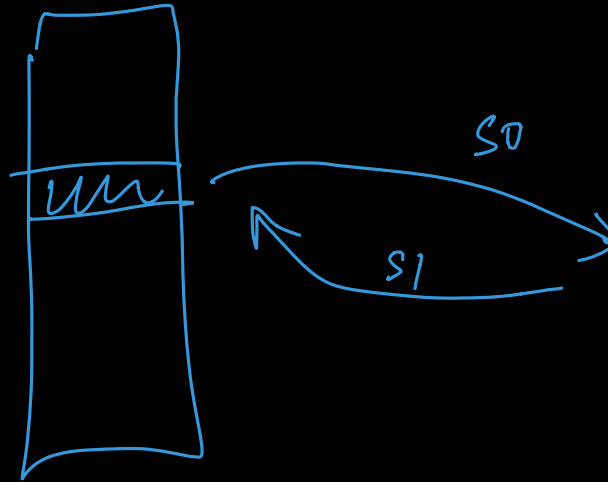
- A process can run without having all of its memory allocated
 - It's allocated on demand
- If the $\{\text{address space used by all processes} + \text{OS}\} \leq \text{physical memory}$ then we're ok
- Otherwise:
 - Make room: ~~discard~~ or store a page onto the disk
 - If the page came from a file & was not modified
 - Discard ... we can always get it
 - If the page is dirty it must be saved in a swap file
 - Swap file: a file (or disk partition) that holds excess pages

Cost

- Handle page fault exception: ~ 400 usec
- Disk seek & read: ~ 10 msec
- Memory access: ~ 100 ns
- Page fault degrades performance by around 100,000!!
- Avoid page faults!
 - If we want $< 10\%$ degradation of performance, we can have just one page fault per 1,000,000 memory accesses

Page replacement

- We need a good replacement policy for good performance



FIFO Replacement

- First In, First Out
- Good
 - May get rid of initialization code or other code that's no longer used
- Bad
 - May get rid of a page holding frequently used global variables

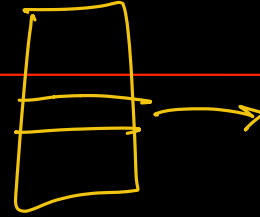
Least Recently Used (LRU)

- Timestamp a page when it is accessed
- When we need to remove a page, search for the one with the oldest timestamp
- Nice algorithm but...
 - Timestamping is a pain – we can't do it with the MMU!

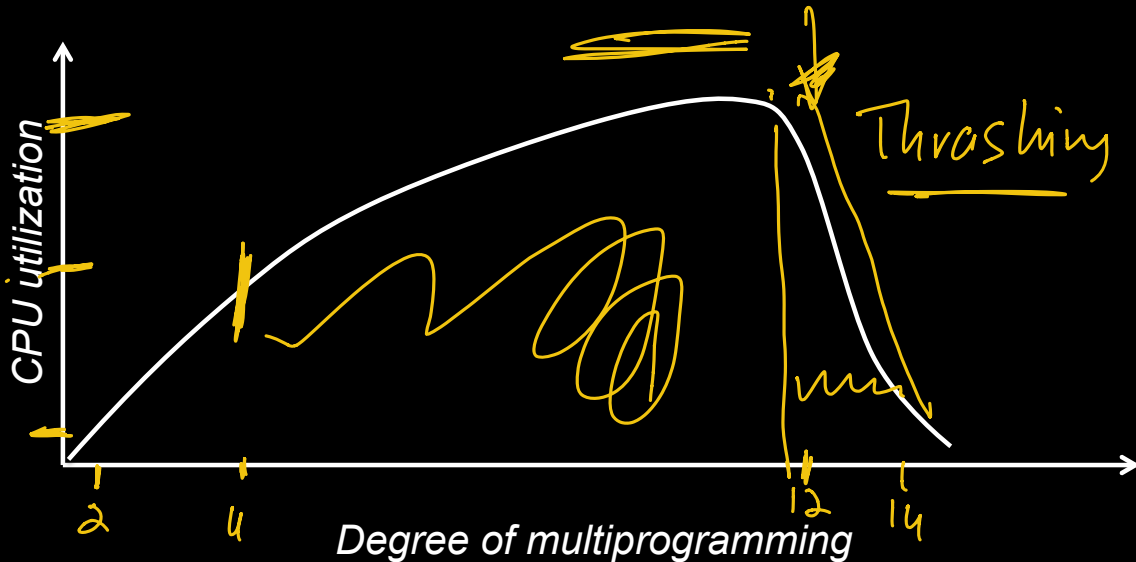
Not Frequently Used Replacement

- ~~Approximate LRU~~
- Each PTE has a reference bit
- Keep a counter for each page frame
- At each clock interrupt:
 - Add the reference bit of each frame to its counter
 - Clear reference bit
- To evict a page, choose the frame with the lowest counter
- Problem
 - No sense of time: a page that was used a lot a long time ago may still have a high count
 - Updating counters is expensive

Thrashing



- Locality
 - Process migrates from one locality (working set) to another
- Thrashing
 - Occurs when sum of all working sets $>$ total memory



Cache Coherence

loading
expected
data
into cache

Faster



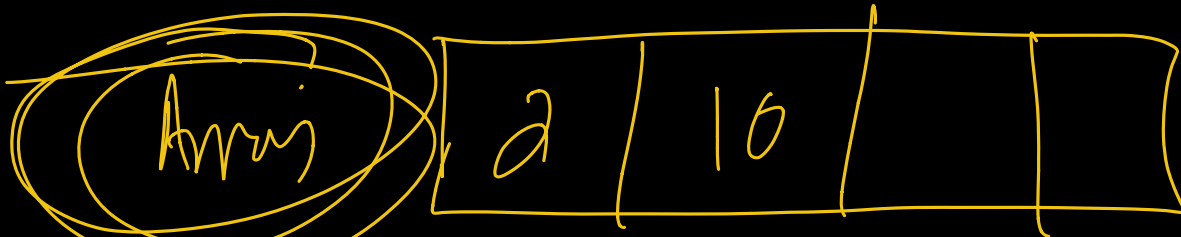
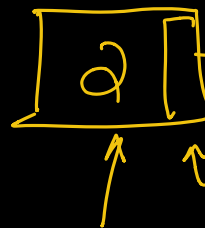
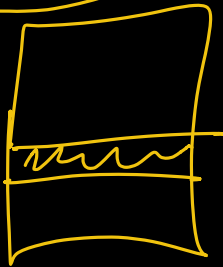
Page Repl.

RAM

TLB etc.

CACHE

CPU



The End